

# Predicting Credit Card Customer Churn with Data in R

*Janhvi Goje*

## Introduction

### Purpose

Our purpose with this project is to **predict if a customer is going to close their credit card account**, so that we can use this information to maximise customer retention by introducing more convenient services in order to convince the customer against leaving the bank.

### Data Description

First, let's start off by looking at all the data that is provided to us and how we can use them in our project.

- **Bank Account Data:** This dataset provides a comprehensive view of customer interactions, account details, and financial behaviors over a period, which are invaluable for predicting customer churn, specifically in terms of closing credit card accounts.
- **Purchase Data:** This dataset offers a rich set of quantitative variables reflecting customers' annual expenditures across various product categories. This data will be analysed to identify homogeneous groups of customers with similar purchasing behaviors.

## Part 1 - Classification

Let's import our `bank_accounts_train.csv` dataset.

```
options(warn = -1)
bank_accounts <- read.csv("/Users/janvigoje/Desktop/bank_accounts_train.csv", colClasses = "character")
head(bank_accounts)
```

```
## CLIENTNUM Customer_Age Gender Dependent_count Education_Level Marital_Status
## 1 721038408          38      F              2      High School      Married
## 2 711968358          55      F              4      Graduate        Single
## 3 719174433          36      F              1      Unknown          Single
## 4 712841883          50      M              3      College           Single
## 5 715980258          36      F              1      Doctorate         Single
## 6 708394908          44      M              0      High School        Single
## Card_Category Months_on_book Total_Relationship_Count Months_Inactive_12_mon
## 1      Blue          27              4              1
## 2      Blue          44              1              2
## 3      Blue          24              2              1
## 4      Silver        36              3              3
## 5      Blue          25              1              2
## 6      Blue          33              5              1
## Contacts_Count_12_mon Credit_Limit Total_Revolving_Bal Avg_Open_To_Buy
## 1          2          1830              0          1830
## 2          1          2638          1423          1215
```

```
## 3          2          1735          0          1735
## 4          2          34516          2517          31999
## 5          3          2835          2511          324
## 6          3          2337          1442          895
##   Total_Amt_Chng_Q4_Q1 Total_Trans_Amt Total_Trans_Ct Total_Ct_Chng_Q4_Q1
## 1          736          1741          43          654
## 2          551          4153          77          925
## 3          74          2467          35          346
## 4          735          7155          66          65
## 5          731          4811          77          925
## 6          659          3945          65          548
##   Avg_Utilization_Ratio          Income Closed_Account
## 1          0 169.7796677049005          0
## 2          539 87.02969695441425          0
## 3          0 140.2794818393886          1
## 4          73 90.16357845626771          1
## 5          886 101.55167812481524          0
## 6          617 46.926825870759785          0
```

## Data Cleaning (Q1)

We are informed that the categorical features have some missing values denoted by "Unknown". Instinctively, we can simply remove these rows; but this might lead to loss of information. Therefore, we replace Unknown with 'NA'.

```
cols <- colnames(bank_accounts)
Data <- bank_accounts          # Duplicate data frame
Data[Data == "Unknown"] <- NA # Replace particular value with NA
```

Next, we turn our categorical variables into factors.

```
categ = c("Education_Level", "Marital_Status", "Card_Category", "Gender", "Closed_Account")
for(i in 1:ncol(Data)){
  if (cols[i] %in% categ){
    Data[,i] = as.factor(Data[,i])
  }
}
```

Next, we turn the strings that are numerical, into numeric.

```
for(i in 1:ncol(Data)){
  if (!(cols[i] %in% categ)){
    Data[,i] = as.numeric(Data[,i])
  }
}
```

```
head(Data)
```

```
##   CLIENTNUM Customer_Age Gender Dependent_count Education_Level Marital_Status
## 1 721038408          38     F           2      High School      Married
## 2 711968358          55     F           4      Graduate      Single
## 3 719174433          36     F           1      <NA>      Single
## 4 712841883          50     M           3      College      Single
## 5 715980258          36     F           1      Doctorate      Single
## 6 708394908          44     M           0      High School      Single
##   Card_Category Months_on_book Total_Relationship_Count Months_Inactive_12_mon
## 1      Blue          27              4              1
```

## 2	Blue	44	1	2
## 3	Blue	24	2	1
## 4	Silver	36	3	3
## 5	Blue	25	1	2
## 6	Blue	33	5	1
##	Contacts_Count_12_mon	Credit_Limit	Total_Revolving_Bal	Avg_Open_To_Buy
## 1	2	1830	0	1830
## 2	1	2638	1423	1215
## 3	2	1735	0	1735
## 4	2	34516	2517	31999
## 5	3	2835	2511	324
## 6	3	2337	1442	895
##	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1
## 1	736	1741	43	654
## 2	551	4153	77	925
## 3	74	2467	35	346
## 4	735	7155	66	65
## 5	731	4811	77	925
## 6	659	3945	65	548
##	Avg_Utilization_Ratio	Income	Closed_Account	
## 1	0	169.77967	0	
## 2	539	87.02970	0	
## 3	0	140.27948	1	
## 4	73	90.16358	1	
## 5	886	101.55168	0	
## 6	617	46.92683	0	

## Exploratory Data Analysis (Q2)

Let's start with our EDA, looking at the columns we have available.

```
colnames(Data)
```

```
## [1] "CLIENTNUM"          "Customer_Age"
## [3] "Gender"              "Dependent_count"
## [5] "Education_Level"     "Marital_Status"
## [7] "Card_Category"       "Months_on_book"
## [9] "Total_Relationship_Count" "Months_Inactive_12_mon"
## [11] "Contacts_Count_12_mon" "Credit_Limit"
## [13] "Total_Revolving_Bal"  "Avg_Open_To_Buy"
## [15] "Total_Amt_Chng_Q4_Q1" "Total_Trans_Amt"
## [17] "Total_Trans_Ct"       "Total_Ct_Chng_Q4_Q1"
## [19] "Avg_Utilization_Ratio" "Income"
## [21] "Closed_Account"
```

Next, let's make some visualizations. We start by making histograms for all the continuous variables. Note that here, the red line is the median and the black line is the mean.

```
contcols <- c("Customer_Age", "Credit_Limit", "Total_Revolving_Bal",
              "Avg_Open_To_Buy", "Total_Amt_Chng_Q4_Q1", "Total_Trans_Amt",
              "Total_Trans_Ct", "Total_Ct_Chng_Q4_Q1", "Avg_Utilization_Ratio",
              "Income", "Months_on_book")

for (col in colnames(Data)) {
  if (col %in% contcols) {
    print(" ")
  }
}
```

```

hist(Data[[col]],
      main = paste("Histogram of", col),
      xlab = col,
      breaks = 20,
      freq = TRUE,
      border = 0,
      col = rgb(0, 0, 0.6, 0.4),
      ann = FALSE,
      axes = FALSE)

axis(1, cex.axis = 0.9, at = seq(min(Data[[col]], na.rm = TRUE), max(Data[[col]], na.rm = TRUE), length.out = 5), las = 2, cex.axis = 0.9)


mtext(col, side = 3, cex = 2)

abline(v = mean(Data[[col]], na.rm = TRUE), lwd = 3)
text(mean(Data[[col]], na.rm = TRUE), max(0.7 * max(table(cut(Data[[col]], breaks = 20))), na.rm = TRUE), pos = 1)

abline(v = median(Data[[col]], na.rm = TRUE), lwd = 3, col = 2)
text(median(Data[[col]], na.rm = TRUE), max(0.7 * max(table(cut(Data[[col]], breaks = 20))), na.rm = TRUE), pos = 1)
}
}


```

```
## [1] " "
```



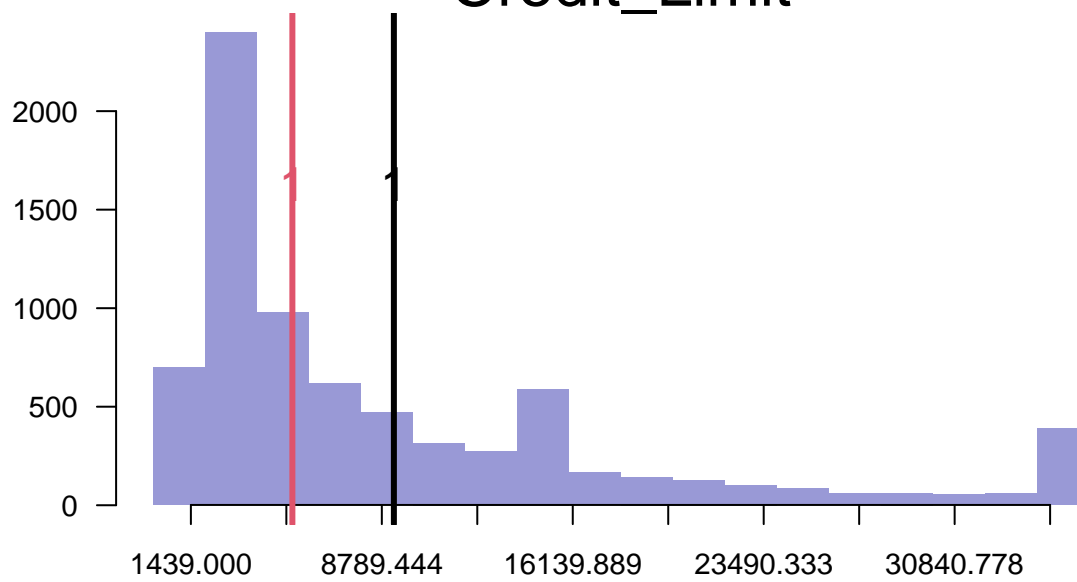
```
## [1] " "
```

## Months\_on\_book

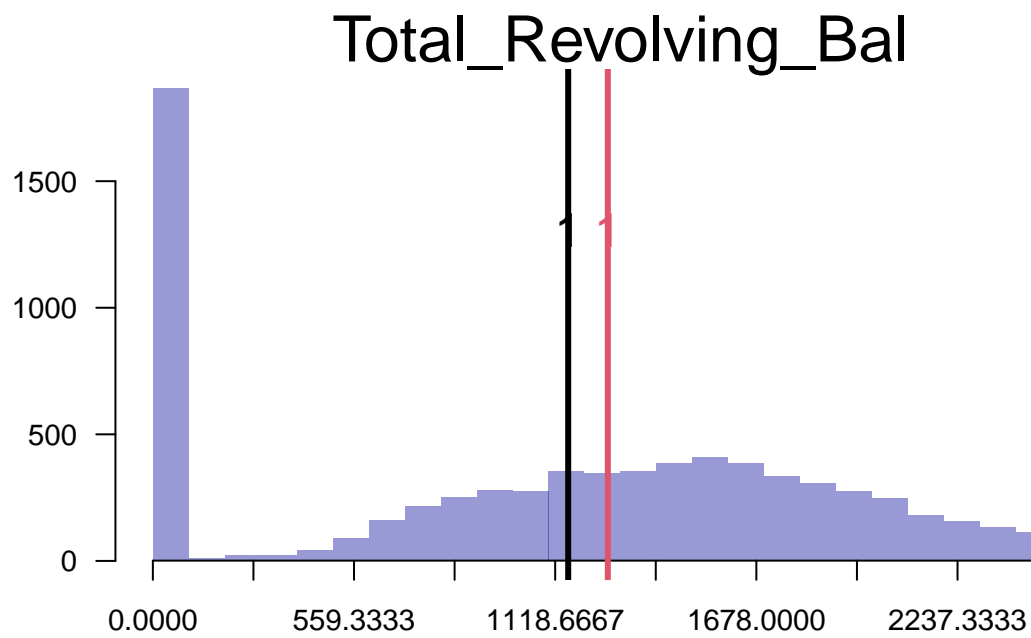


## [1] " "


## Credit\_Limit



## [1] " "




# Total\_Amt\_Chng\_Q4\_Q1

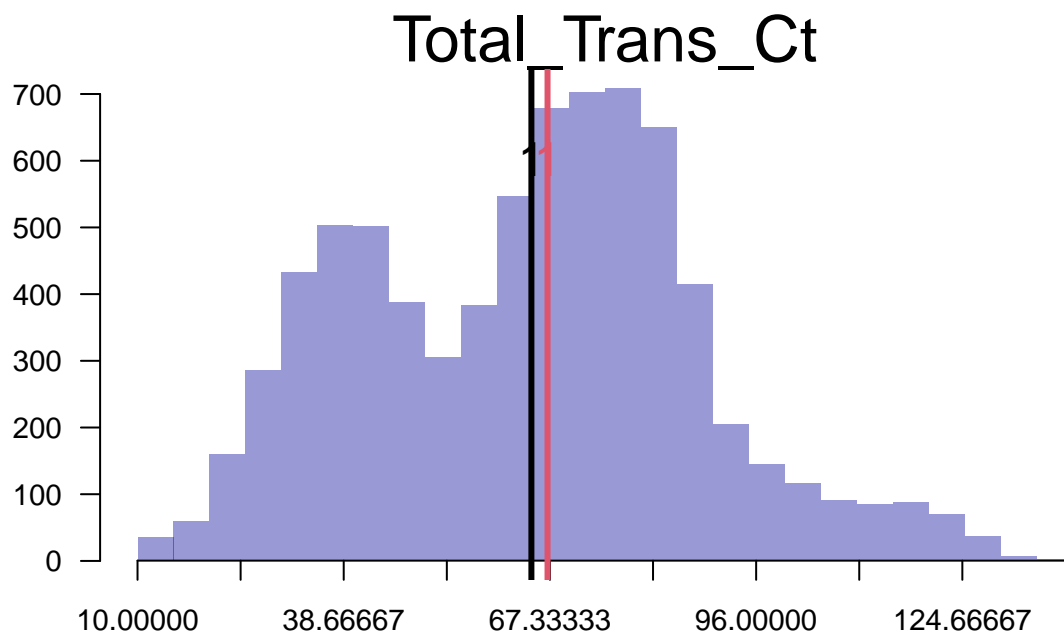


## [1] " "

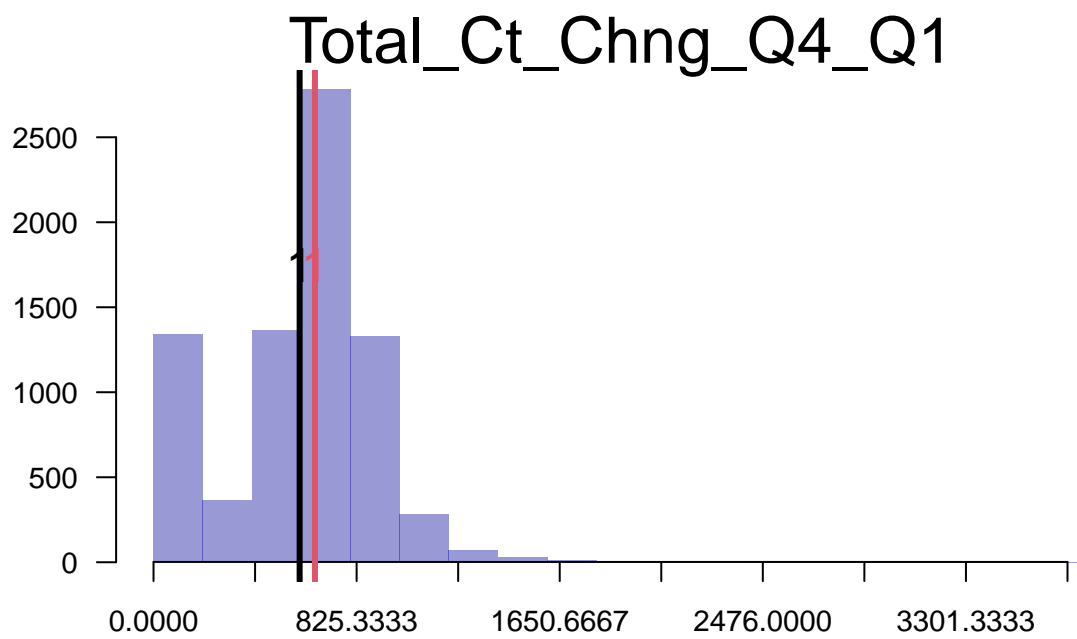
# Total\_Trans\_Amt



## [1] " "

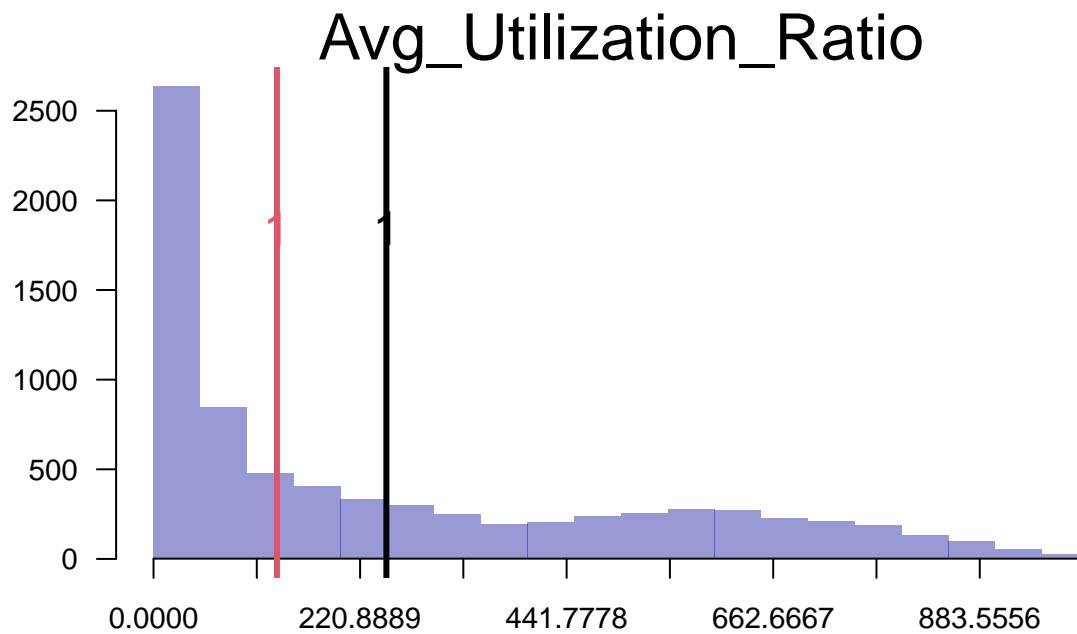


## [1] " "

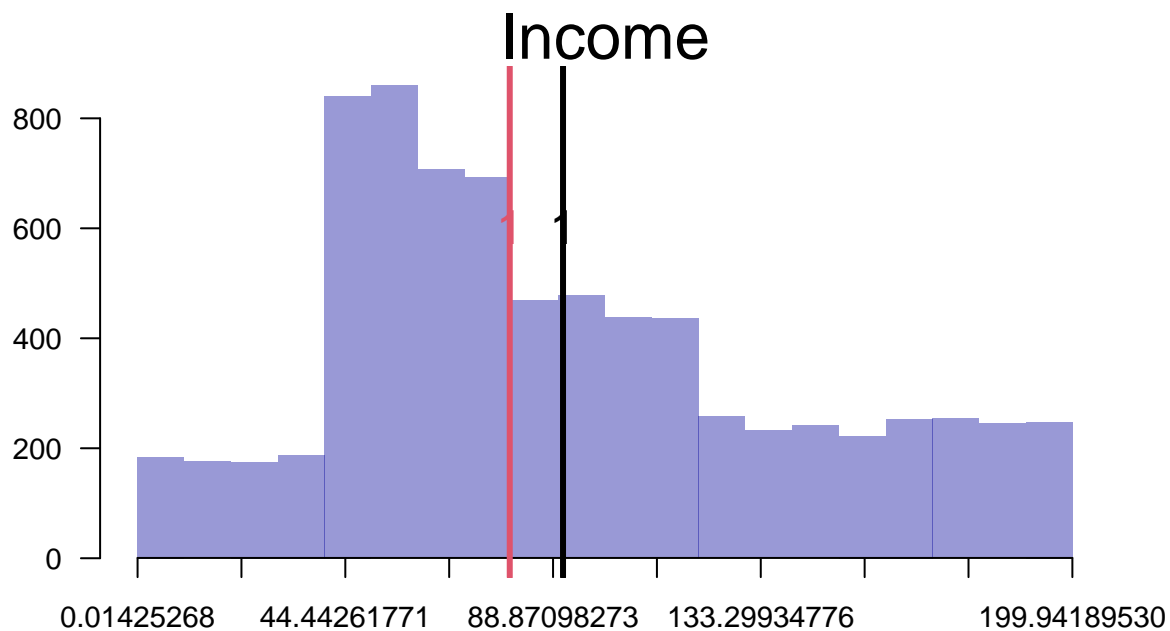


## [1] " "





```
## [1] " "
```



The histograms for 'Total\_Revolving\_Bal', 'Avg\_Open\_To\_Buy', 'Total\_Trans\_Amt', 'Avg\_Utilization\_Ratio', and 'Income' exhibit right-skewed distributions, indicating that most customers have lower balances, credit available, transaction amounts, utilization ratios, and incomes, with fewer customers having higher values in these areas. In contrast, 'Total\_Trans\_Ct' shows a roughly normal distribution, suggesting consistent transaction behavior across customers. The 'Total\_Amt\_Chng\_Q4\_Q1' histogram also appears symmetrical, centering around the mean, which indicates uniform changes in transaction amounts over time.

Let's look at our categorical variables now. A pie chart is ideal for categories like marital status and gender with low number of modalities.

```
pivars <- c("Gender", "Education_Level", "Marital_Status", "Card_Category", "Closed_Account")
```

```

for (var in pivars) {
  print(" ")
  tab <- table(Data[[var]])

  pct <- round((tab / sum(tab)) * 100, 1)

  labels <- paste(pct, "%", sep = "")
  legend_labels <- paste(names(tab), ":", labels)


  pie(x = tab,
      labels = NA,
      main = paste("Distribution of", var),
      col = rainbow(length(tab)),
      cex = 1.2,
      init.angle = 90,
      radius = 1)

  legend("topright",
        legend = legend_labels,
        pch = 15,
        col = rainbow(length(tab)),
        cex = 0.8,
        bty = "n")
}

```


```
## [1] " "
```

## Distribution of Gender



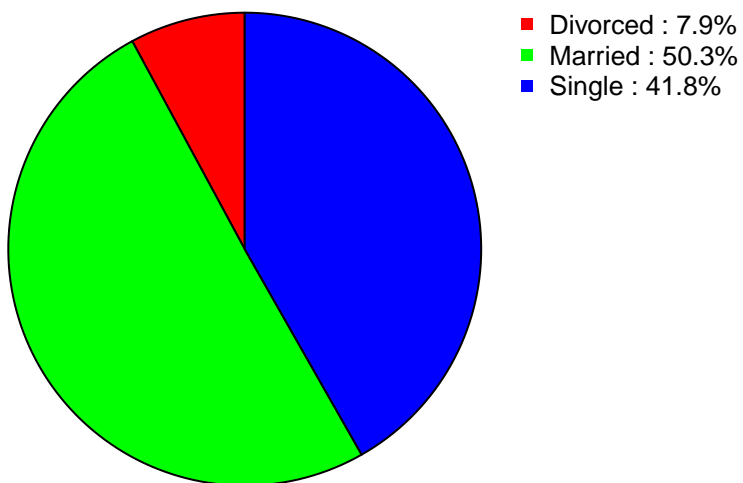
```
## [1] " "
```

## Distribution of Education\_Level



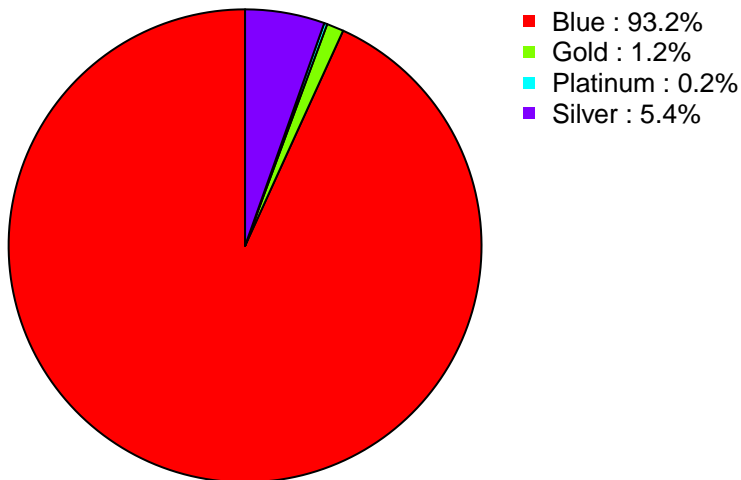
## [1] " "

## Distribution of Marital\_Status




## [1] " "

## Distribution of Card\_Category



```
## [1] " "
```

## Distribution of Closed\_Account



The 'Distribution of Gender' chart shows a nearly even split between male and female, with females slightly more prevalent. The 'Distribution of Education\_Level' indicates a diverse educational background among the individuals, with the majority holding a graduate education. The 'Distribution of Marital\_Status' reveals that a significant proportion of individuals are married. The 'Distribution of Card\_Category' shows that most individuals have a Blue card. Lastly, the 'Distribution of Closed\_Account' indicates that the majority of accounts are active.

Now, let's visualise columns with high modalities with a Bar Plot.

```
barvars <- c("Dependent_count", "Total_Relationship_Count",  
             "Months_Inactive_12_mon", "Contacts_Count_12_mon")  
for (var in barvars) {  
  print(" ")  
  var_table <- table(Data[[var]])
```

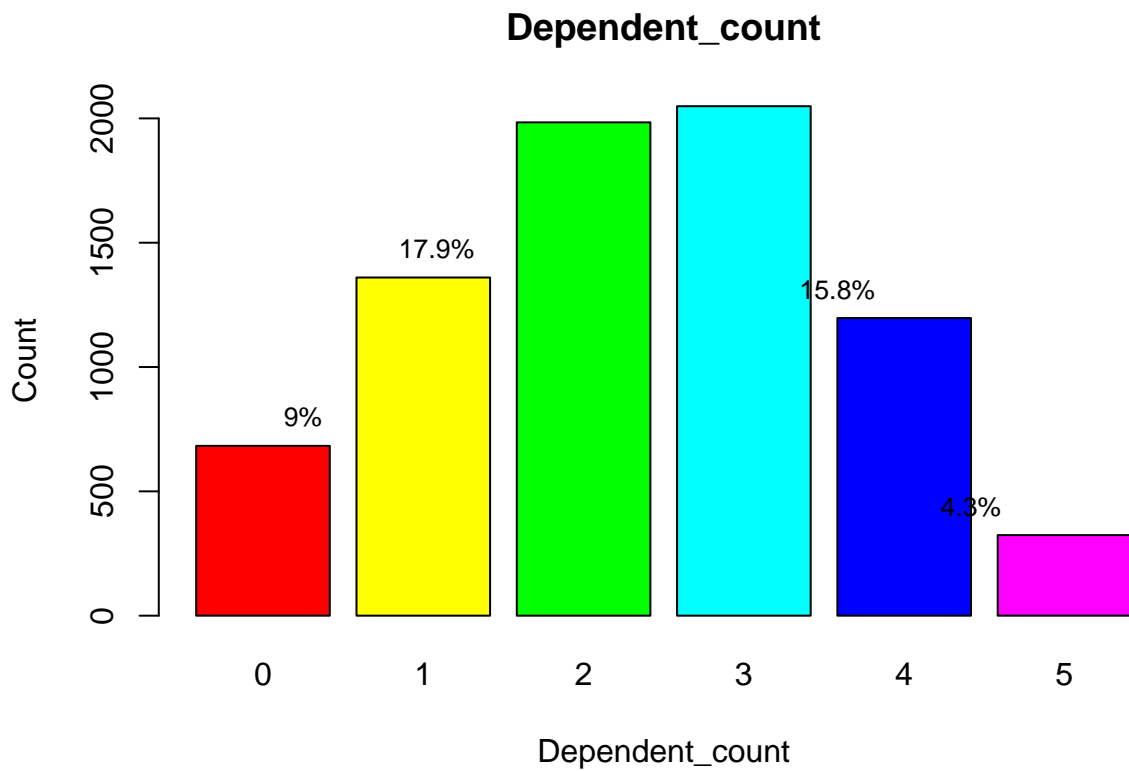
```
percentages <- round((var_table / sum(var_table)) * 100, 1)

barplot(var_table,
        col = rainbow(length(var_table)), # Use rainbow colors for differentiation
        main = var, # Use variable name as the main title
        ylim = c(0, max(var_table) + 5), # Adjust y limit for space for labels
        xlab = var, # Label x-axis with variable name
        ylab = "Count") # Label y-axis as "Count"

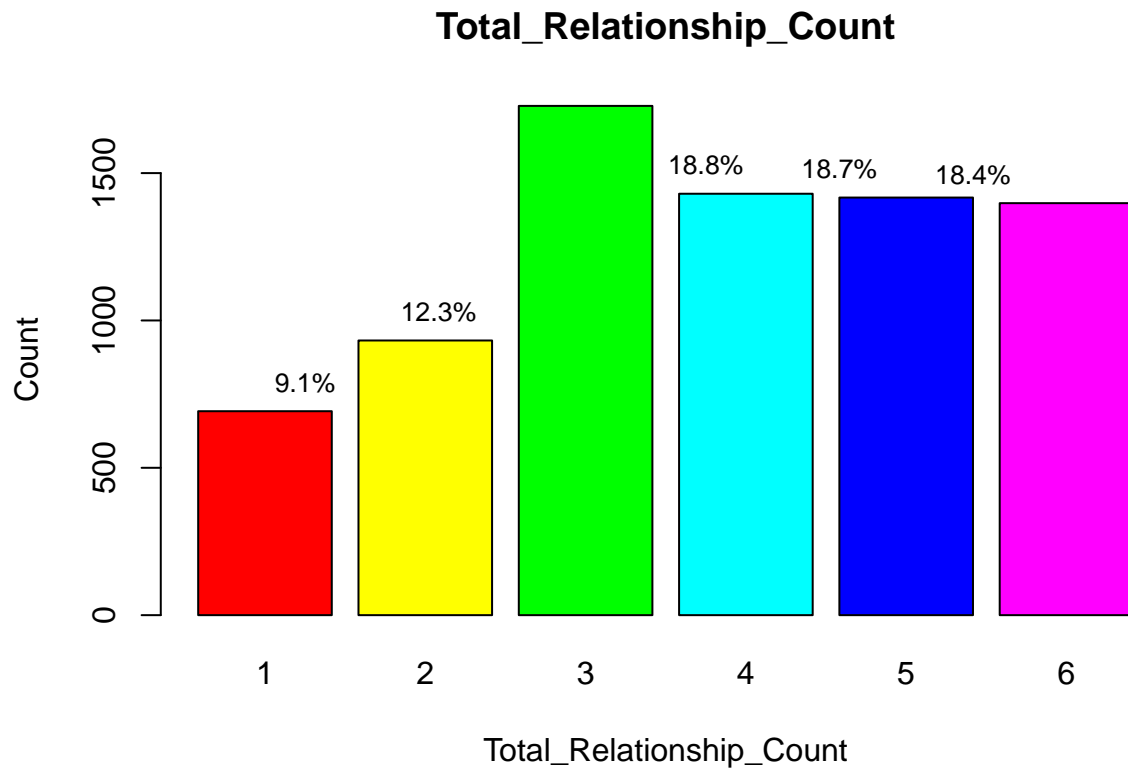
text(x = seq_along(var_table), y = var_table + 1, labels = paste(percentages, "%", sep=""), pos = 3, c
}

```

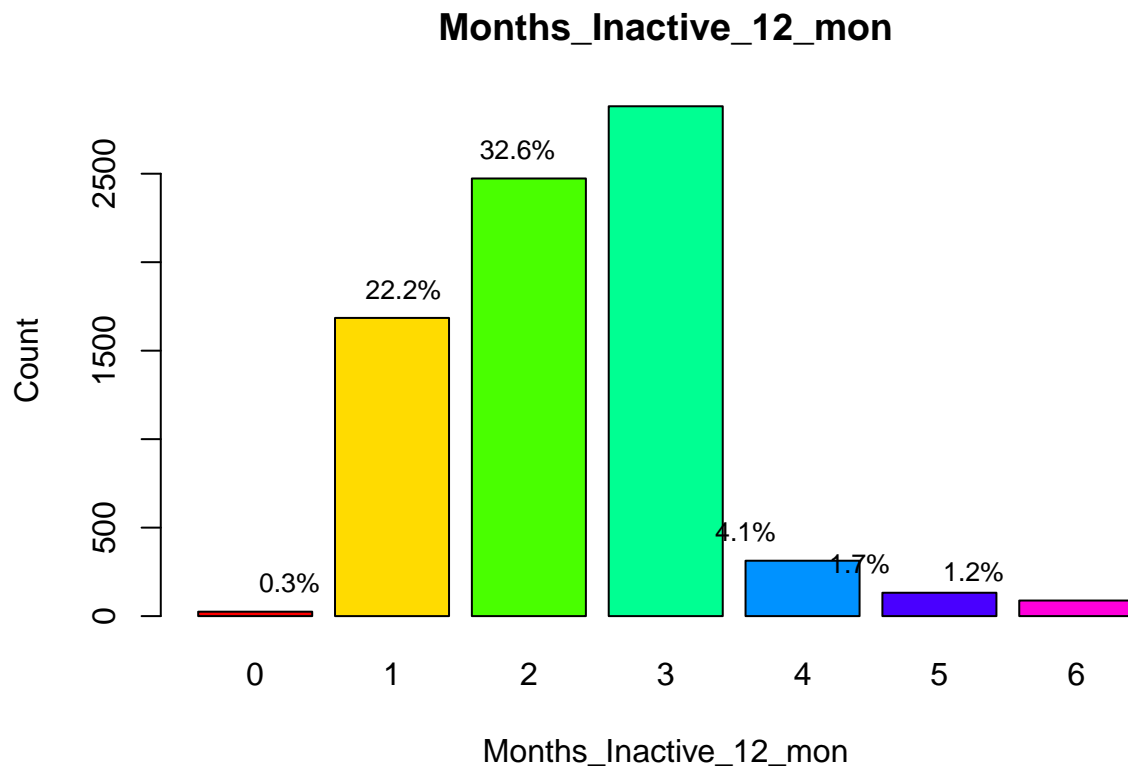
```
## [1] " "
```



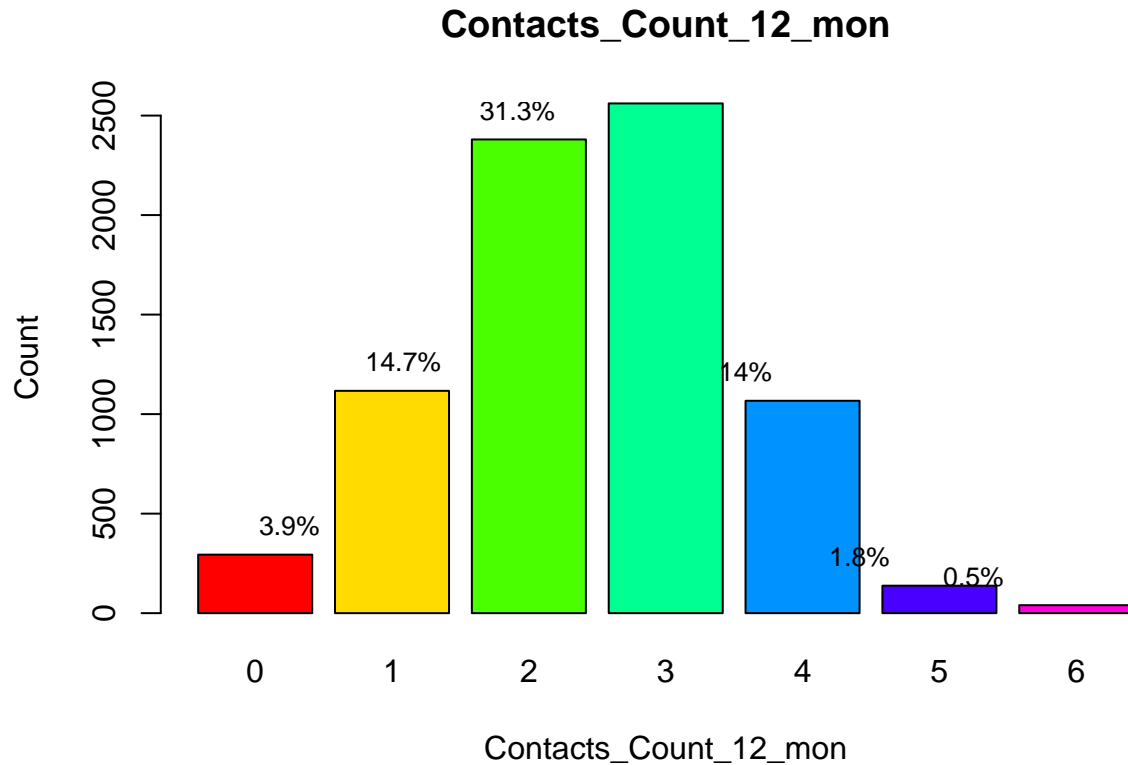
```
## [1] " "
```



## [1] " "



## [1] " "




These visualisation are great to give us a general overlook of our data. We can make some interesting observations:

Dependent\_count: The chart indicates that a considerable number of customers have two or three dependents. Total\_Relationship\_Count: Customers are spread fairly evenly across different relationship levels with the bank. Months\_Inactive\_12\_mon: There is a significant concentration of customers who have been inactive for three months within the last year. Contacts\_Count\_12\_mon: The distribution suggests that contacting customers two or three times in the past 12 months is most common

## Bivariate Analysis

Now, let's move on to looking at relevant relationships between variables. Let's look at a few relevant categorical relationships with a closed account. Closed account remains to be the centre of our focus since we keep our purpose (defined in the introduction) in mind.


```
library(ggplot2)
ggplot(Data, aes(x = Gender, fill = factor(Closed_Account))) +
  geom_bar(position = "fill") +
  scale_y_continuous(labels = scales::percent_format()) +
  labs(x = "Gender", y = "Percentage", fill = "Closed Account") +
  theme_minimal()
```



The chart shows a similar small proportion of closed accounts for both genders, with the vast majority of accounts remaining active for both females and males.


```
ggplot(Data, aes(x = Education_Level, fill = factor(Closed_Account))) +  
  geom_bar(position = "fill") +  
  scale_y_continuous(labels = scales::percent_format()) +  
  labs(x = "Education Level", y = "Percentage", fill = "Closed Account") +  
  theme_minimal()
```





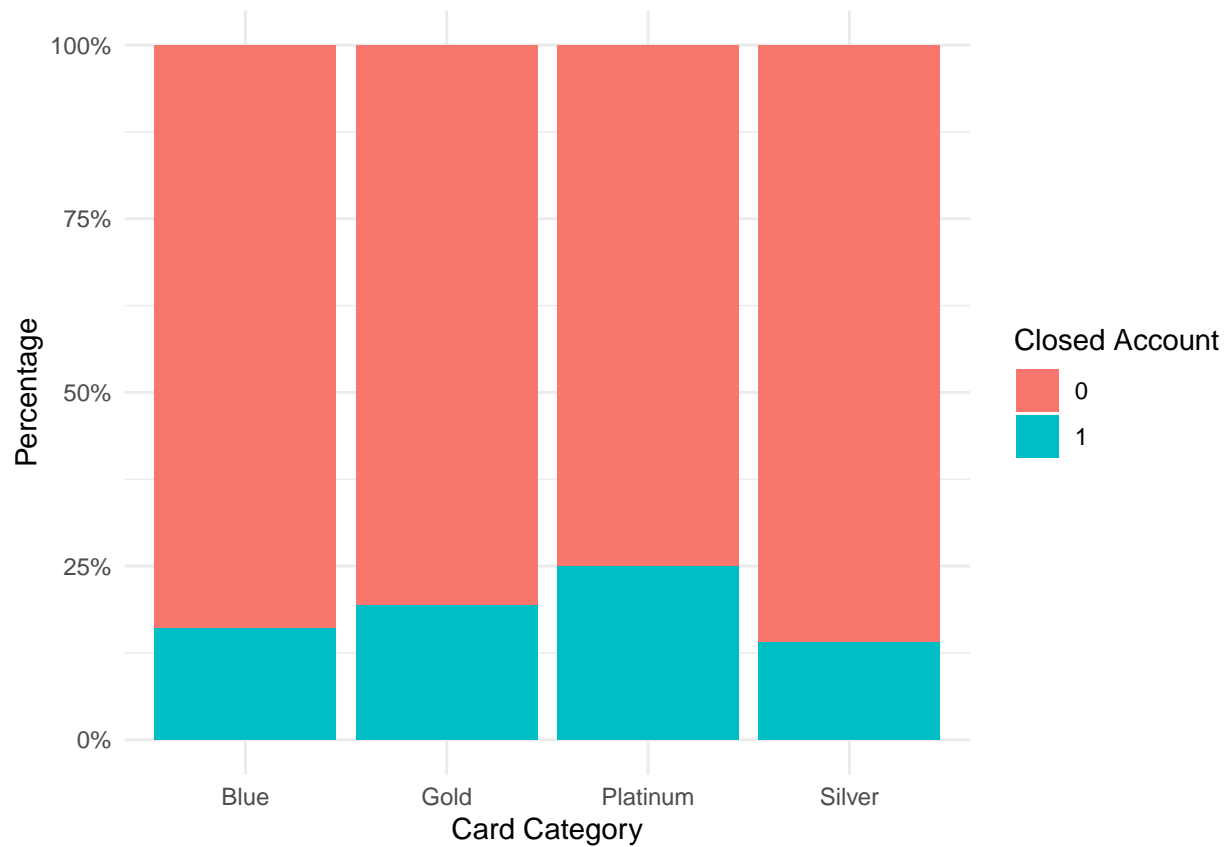
Across different education levels, the majority of accounts are active, with closure rates being comparatively low.

```
ggplot(Data, aes(x = Marital_Status, fill = factor(Closed_Account))) +
  geom_bar(position = "fill") +
  scale_y_continuous(labels = scales::percent_format()) +
  labs(x = "Marital Status", y = "Percentage", fill = "Closed Account") +
  theme_minimal()
```




Active accounts predominate across all marital statuses, with closed accounts making up a smaller portion.

```
ggplot(Data, aes(x = Card_Category, fill = factor(Closed_Account))) +  
  geom_bar(position = "fill") +  
  scale_y_continuous(labels = scales::percent_format()) +  
  labs(x = "Card Category", y = "Percentage", fill = "Closed Account") +  
  theme_minimal()
```



The blue card category shows the most activity with very few closures, consistent across other card types.

```
ggplot(Data, aes(x = Total_Relationship_Count, fill = factor(Closed_Account))) +  
  geom_bar(position = "fill") +  
  scale_y_continuous(labels = scales::percent_format()) +  
  labs(x = "Total Relationship Count", y = "Percentage", fill = "Closed Account") +  
  theme_minimal()
```



Most customers maintain active accounts across all levels of product relationships with the bank.

Let's perform some hypothesis testing to rigidify our interpretations.

```
# Ensure the Data$Closed_Account is a factor for accurate Chi-square testing
Data$Closed_Account <- factor(Data$Closed_Account)
```

```
# Gender vs. Closed_Account
# H0: There is no association between Gender and Closed_Account status.
# Ha: There is an association between Gender and Closed_Account status.
gender_table <- table(Data$Gender, Data$Closed_Account)
chisq.test(gender_table)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data: gender_table
## X-squared = 12.231, df = 1, p-value = 0.0004699
```

With a p-value of 0.0004699, there is strong evidence against the null hypothesis. This suggests there is an association between gender and whether an account is closed.

```
# Education_Level vs. Closed_Account
# H0: There is no association between Education_Level and Closed_Account status.
# Ha: There is an association between Education_Level and Closed_Account status.
education_level_table <- table(Data$Education_Level, Data$Closed_Account)
chisq.test(education_level_table)
```

```
##
```

```
## Pearson's Chi-squared test
##
## data:  education_level_table
## X-squared = 12.339, df = 5, p-value = 0.03043
```

The p-value of 0.03043 indicates that there is a statistically significant association between education level and closed account status, although the evidence is not as strong as for gender.

```
# Marital_Status vs. Closed_Account
# H0: There is no association between Marital_Status and Closed_Account status.
# Ha: There is an association between Marital_Status and Closed_Account status.
marital_status_table <- table(Data$Marital_Status, Data$Closed_Account)
chisq.test(marital_status_table)
```

```
##
## Pearson's Chi-squared test
##
## data:  marital_status_table
## X-squared = 6.6217, df = 2, p-value = 0.03648
```

The p-value of 0.03648 suggests a statistically significant association between marital status and closed account status, but this significance is marginal and could be sensitive to sample size and distribution.

```
# Card_Category vs. Closed_Account
# H0: There is no association between Card_Category and Closed_Account status.
# Ha: There is an association between Card_Category and Closed_Account status.
card_category_table <- table(Data$Card_Category, Data$Closed_Account)
chisq.test(card_category_table)
```

```
##
## Pearson's Chi-squared test
##
## data:  card_category_table
## X-squared = 2.8848, df = 3, p-value = 0.4097
```

With a p-value of 0.4097, there is insufficient evidence to suggest an association between card category and closed account status, and we fail to reject the null hypothesis.

```
# Total_Relationship_Count vs. Closed_Account
# H0: There is no association between Total_Relationship_Count and Closed_Account status.
# Ha: There is an association between Total_Relationship_Count and Closed_Account status.
total_relationship_count_table <- table(Data$Total_Relationship_Count, Data$Closed_Account)
chisq.test(total_relationship_count_table)
```

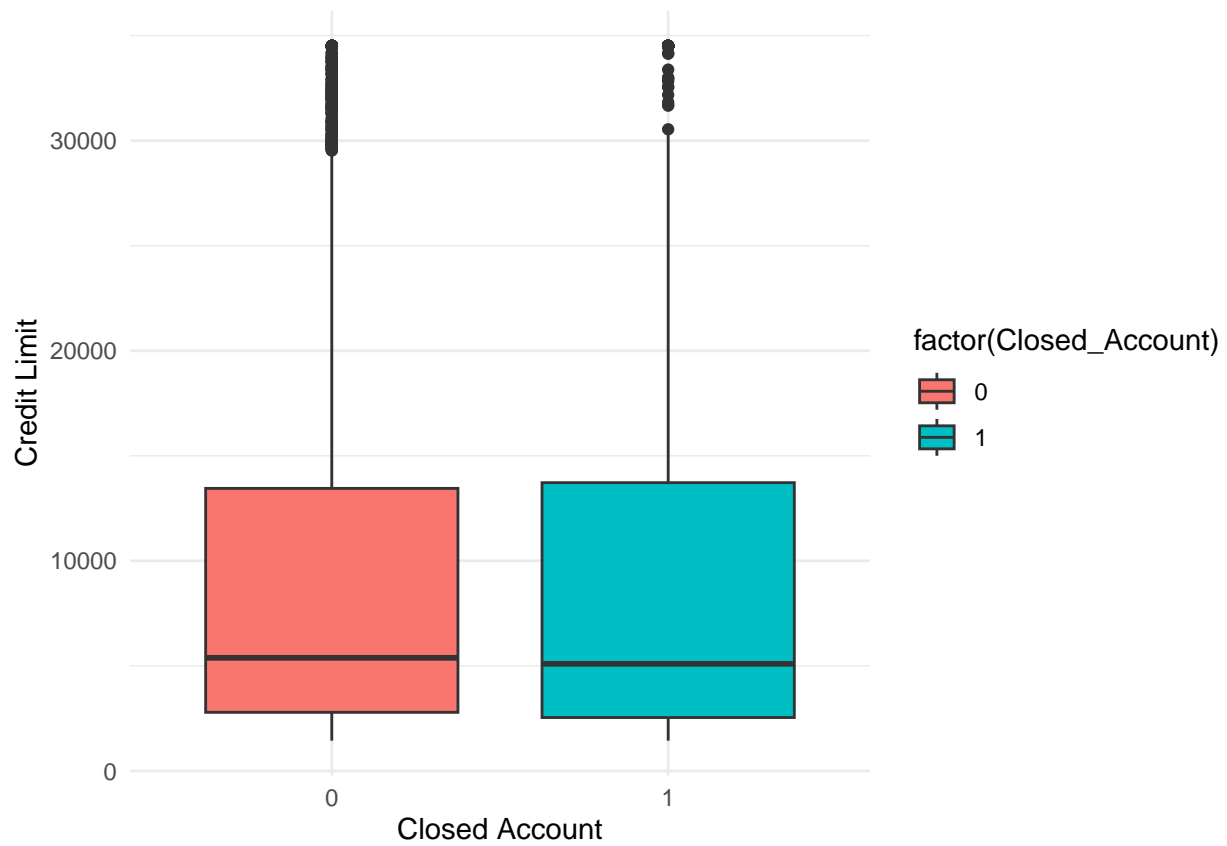
```
##
## Pearson's Chi-squared test
##
## data:  total_relationship_count_table
## X-squared = 215.5, df = 5, p-value < 2.2e-16
```

The very small p-value ( $< 2.2e-16$ ) indicates a very strong association between the total relationship count and closed account status, strongly rejecting the null hypothesis.

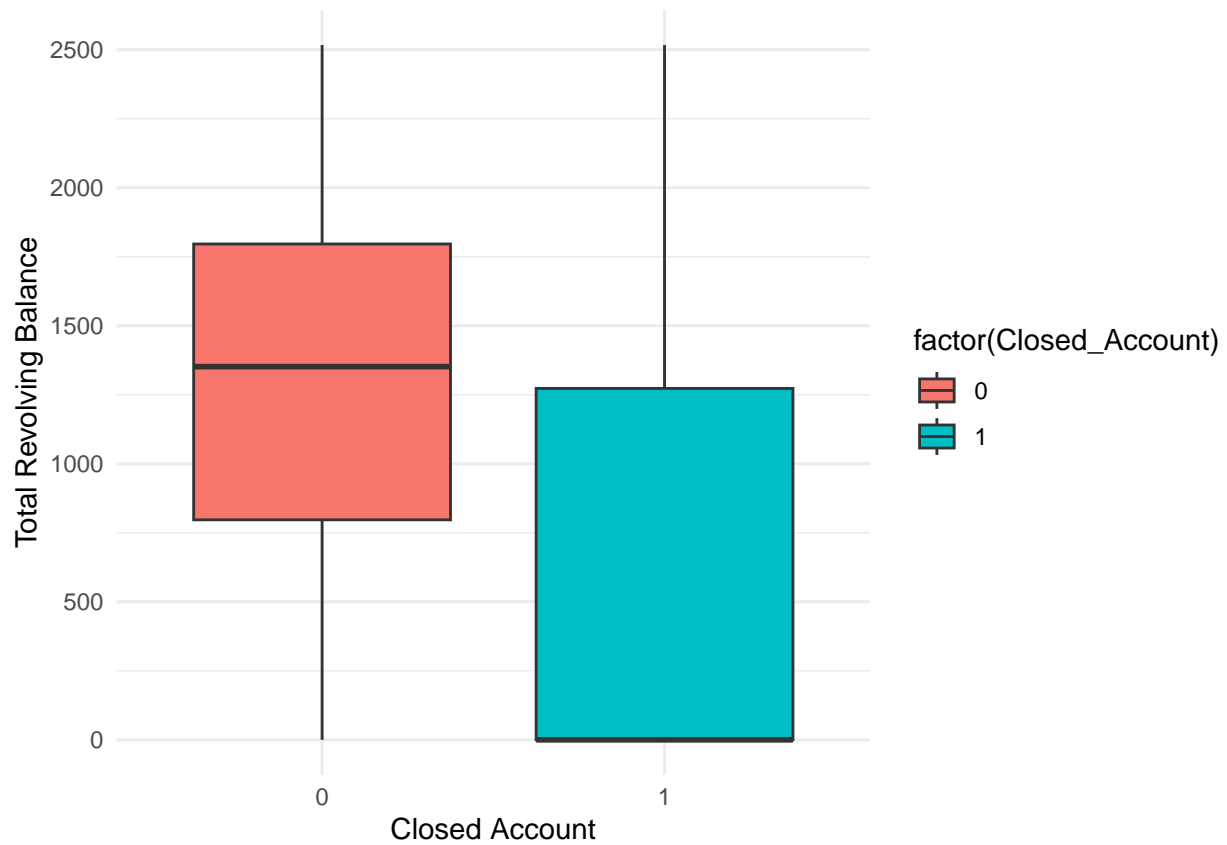
Now, let's do a boxplot to understand the relationship between Closed\_Account and the relevant continuous variables.

```
ggplot(Data, aes(x=factor(Closed_Account), y=Credit_Limit, fill=factor(Closed_Account))) +
  geom_boxplot() +
  labs(x="Closed Account", y="Credit Limit") +
```


```
theme_minimal()
```



```
ggplot(Data, aes(x=factor(Closed_Account), y=Total_Revolving_Bal, fill=factor(Closed_Account))) +  
  geom_boxplot() +  
  labs(x="Closed Account", y="Total Revolving Balance") +  
  theme_minimal()
```




```
ggplot(Data, aes(x=factor(Closed_Account), y=Customer_Age, fill=factor(Closed_Account))) +  
  geom_boxplot() +  
  labs(x="Closed Account", y="Customer Age") +  
  theme_minimal()
```




```
ggplot(Data, aes(x=factor(Closed_Account), y=Income, fill=factor(Closed_Account))) +  
  geom_boxplot() +  
  labs(x="Closed Account", y="Income") +  
  theme_minimal()
```





```
ggplot(Data, aes(x=factor(Closed_Account), y=Avg_Utilization_Ratio, fill=factor(Closed_Account))) +  
  geom_boxplot() +  
  labs(x="Closed Account", y="Average Utilization Ratio") +  
  theme_minimal()
```



The boxplots reveal that active accounts generally have higher credit limits and revolving balances than closed accounts, with a broader range in both metrics. Customer age and income levels are comparable for both active and closed accounts, suggesting that account status does not vary significantly with these demographics. However, active accounts exhibit a higher median credit utilization ratio, with a wider spread indicating greater variability in how much credit is used relative to the limit, compared to closed accounts which show a more condensed utilization pattern.

Now, let's perform t-test to determine further the impact of these variables on the Closed\_Account variable.

```
t_test_credit_limit <- t.test(Data$Credit_Limit ~ Data$Closed_Account)
print(t_test_credit_limit)
```

```
##
##  Welch Two Sample t-test
##
## data:  Data$Credit_Limit by Data$Closed_Account
## t = 1.121, df = 1727.1, p-value = 0.2625
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
##  -234.7302  860.9402
## sample estimates:
## mean in group 0 mean in group 1
##      9303.600      8990.495
```

With a p-value of 0.2625, there is not enough evidence to conclude that there is a significant difference in credit limits between active and closed accounts. The means of both groups are similar, and the confidence interval includes zero, further suggesting no significant difference.

```
t_test_total_revolving_balance <- t.test(Data$Total_Revolving_Bal ~ Data$Closed_Account)
print(t_test_total_revolving_balance)
```

```
##
## Welch Two Sample t-test
##
## data: Data$Total_Revolving_Bal by Data$Closed_Account
## t = 20.969, df = 1545.5, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## 531.8715 641.6442
## sample estimates:
## mean in group 0 mean in group 1
## 1248.7842 662.0263
```

The p-value is less than 2.2e-16, which is extremely small, indicating a highly significant difference in the total revolving balances between active and closed accounts. The average revolving balance is significantly higher for active accounts than for closed ones.

```
t_test_income <- t.test(Data$Income ~ Data$Closed_Account)
print(t_test_income)
```

```
##
## Welch Two Sample t-test
##
## data: Data$Income by Data$Closed_Account
## t = -0.62389, df = 1689, p-value = 0.5328
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## -4.010509 2.074841
## sample estimates:
## mean in group 0 mean in group 1
## 90.84967 91.81750
```

The p-value of 0.5328 suggests that there is no significant difference in income between active and closed accounts. The means are nearly equal, and the confidence interval includes zero.

```
t_test_avg_utilization_ratio <- t.test(Data$Avg_Utilization_Ratio ~ Data$Closed_Account)
print(t_test_avg_utilization_ratio)
```

```
##
## Welch Two Sample t-test
##
## data: Data$Avg_Utilization_Ratio by Data$Closed_Account
## t = 15.046, df = 1774.8, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## 105.3873 136.9820
## sample estimates:
## mean in group 0 mean in group 1
## 268.4418 147.2572
```

A p-value less than 2.2e-16 signifies a highly significant difference in average utilization ratios between the two groups. Active accounts have a substantially higher average utilization ratio compared to closed accounts.

```
t_test_customer_age <- t.test(Data$Customer_Age ~ Data$Closed_Account)
print(t_test_customer_age)
```

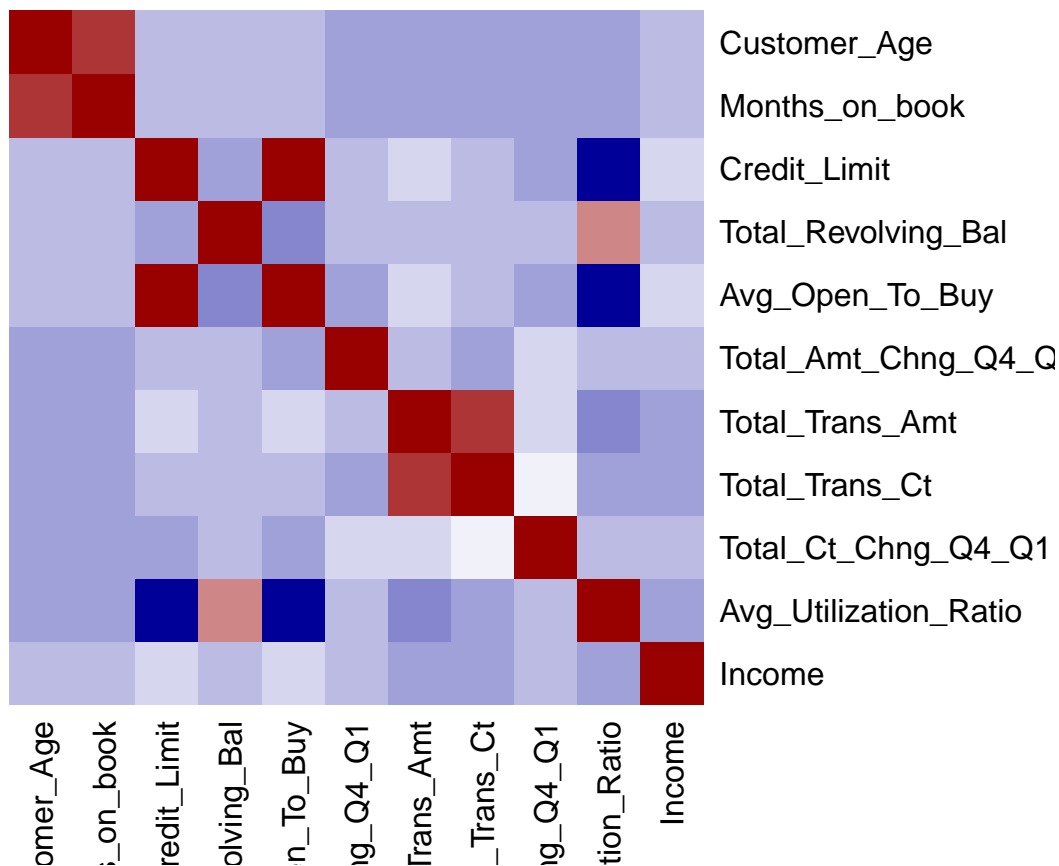
```
##
## Welch Two Sample t-test
##
## data: Data$Customer_Age by Data$Closed_Account
## t = -2.252, df = 1758, p-value = 0.02444
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## -1.02951386 -0.07103968
## sample estimates:
## mean in group 0 mean in group 1
## 46.17774 46.72802
```

With a p-value of 0.02444, there is a statistically significant difference in age between active and closed accounts, although the actual difference in mean ages is quite small.

Finally, let's look into the relationship between certain continuous variables to further gain insight into some interesting patterns that we might find interesting.

```
Data_num <- Data[,names(Data) %in% contcols]
Cor.Matrix = cor(Data_num)

col<- colorRampPalette(c(rgb(0, 0,.6), "white", rgb(.6,0,0)))(20)
heatmap(x = Cor.Matrix, col = col, symm = T, Rowv = NA)
```



```
correlated_var <- matrix(nrow=1, ncol=3)
for(i in 2:ncol(Data_num)){
  if (is.numeric(Data_num[1,i])) {
    for (j in i:ncol(Data_num)) {
```

```

    if (i != j & (Cor.Matrix[i,j] > 0.70 | Cor.Matrix[i,j] < -0.70)){
      correlated_var <- rbind(correlated_var, c(colnames(Data_num)[i],colnames(Data_num)[j],Cor.Matrix[i,j]))
    }
  }
}
}

```

Customer\_Age and Months\_on\_book might have a strong positive correlation, which is logical since older customers are likely to have been with the bank for longer periods. Credit\_Limit seems to have a strong positive correlation with Avg\_Open\_To\_Buy, which makes sense as the open to buy is typically a function of the credit limit minus the current balance. Total\_Trans\_Amt and Total\_Trans\_Ct show a strong positive correlation, suggesting that customers with more transactions tend to also have higher total transaction amounts. Avg\_Utilization\_Ratio seems to have a strong negative correlation with Avg\_Open\_To\_Buy. This would align with the idea that higher utilization ratios (proportion of credit limit used) often mean less credit is available to use. Between credit\_limit and Avg\_Utilization\_Ratio there is a negative relationship. Customers with higher credit limits tend to use a smaller proportion of their available credit.

### Logistic Regression (Q3)

```

X1 <- Data$Gender
X2 <- Data$Income
Y <- as.numeric(Data$Closed_Account) - 1

mod = glm(Y~X1+X2, family = binomial(link = "logit"))
summary(mod)

##
## Call:
## glm(formula = Y ~ X1 + X2, family = binomial(link = "logit"))
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.5952639  0.0709750 -22.476  < 2e-16 ***
## X1M          -0.2232886  0.0632205  -3.532  0.000413 ***
## X2             0.0004261  0.0006311   0.675  0.499513
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 6685.3  on 7596  degrees of freedom
## Residual deviance: 6672.3  on 7594  degrees of freedom
## AIC: 6678.3
##
## Number of Fisher Scoring iterations: 4

```

According to the output we received, we notice that gender is significant while income isn't. We can further analyse output by taking the exponentials and questioning their odds.

```

exp(coefficients(mod)[1])#odds of closing account for F with avg income

## (Intercept)
##      0.202855

```

```
exp((coefficients(mod)[1]) + coefficients(mod)[2]) #odds of closing account for M with avg income
```

```
## (Intercept)
## 0.1622605
```

```
coefficients(mod)[3] #odds ratio of closing when income increases by 1 std deviation
```

```
## X2
## 0.0004261444
```

```
coefficients(mod)[3]/4 #maximal prob increase rate when income goes up
```

```
## X2
## 0.0001065361
```

By looking at the exponentials, we see that the odds of closing, based on an income variation, are fairly low. We can add an interaction term to better interpret the effect of income.

*# Null Hypothesis (H0): The effect of income on the probability of account closure is the same for male.*

*# Alternative Hypothesis (H1): The effect of income on the probability of account closure differs between*

```
mod2 = glm(Y~X1*X2, family = binomial(link = "logit"))
summary(mod2)
```

```
##
## Call:
## glm(formula = Y ~ X1 * X2, family = binomial(link = "logit"))
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.5474767  0.0787725 -19.645  < 2e-16 ***
## X1M          -0.3990676  0.1445507  -2.761  0.00577 **
## X2           -0.0001008  0.0007426  -0.136  0.89207
## X1M:X2        0.0019037  0.0014016   1.358  0.17440
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6685.3  on 7596  degrees of freedom
## Residual deviance: 6670.5  on 7593  degrees of freedom
## AIC: 6678.5
##
## Number of Fisher Scoring iterations: 4
```

We notice that the Interaction term isn't statistically significant. Let's further interpret the coefficients.

```
exp(coefficients(mod2)[3]) #odds of closing when F and income varies by sd
```

```
## X2
## 0.9998992
```

```
exp(coefficients(mod2)[3] + coefficients(mod2)[4]) #odds of closing when M and income varies by sd
```

```
## X2
## 1.001805
```

We get numbers *very close to 1*. This means that there is Income does have an effect, but since the *interaction term isn't significant*, we don't reject the null hypothesis. We now plot without the interaction term.


```
X1 <- as.numeric(X1) - 1
plot(X2[which(X1 == 0)], Y[which(X1 == 0)], col = rgb(.7,0,0,.2), pch = 16)
points(X2[which(X1 == 1)], Y[which(X1 == 1)], col = rgb(0,0,0.7,.2), pch = 16)

f.true0 = function(x) exp(coefficients(mod)[1] + coefficients(mod)[3]*x) / (1 + exp(coefficients(mod)[1]
f.true1 = function(x) exp(coefficients(mod)[1] + coefficients(mod)[3]*x + coefficients(mod)[2]) / (1 + exp(coefficients(mod)[1] + coefficients(mod)[2] + coefficients(mod)[3]*x)

curve(f.true0, add=T, lwd = 2, col = rgb(.8,0,0))
curve(f.true1, add=T, lwd = 2, col = rgb(0,0,0.8))

abline(h = 0.5, col = 3, lwd = 2, lty = 2)

legend("right", c("X2 = 0", "X2 = 1", "True", "Estimated"),
      col = c(rgb(0.8,0,0), rgb(0,0,0.8), 1, 1), pch = c(15,15,NA,NA), pt.cex = c(2,2,NA,NA),
      lty = c(NA,NA, 1,2), lwd = c(NA,NA, 2,2))
```



We notice a very flat estimation. We also notice low probabilities leading to t never predicting a 1. This might be caused by the insignificance of 'Income' variable.

The Confusion Matrix is as follows.

```
predicted.Y = (predict(mod, type = "response") > .5) * 1

True.positive = sum(predicted.Y[which(Y == 1)] == 1)
True.negative = sum(predicted.Y[which(Y == 0)] == 0)
False.positive = sum(predicted.Y[which(Y == 0)] == 1)
False.negative = sum(predicted.Y[which(Y == 1)] == 0)

Confusion.Matrix = matrix(c(True.positive,
                             False.positive,
                             False.negative,
```

```

      True.negative),
      nrow = 2, ncol = 2)

row.names(Confusion.Matrix) = c("Actual Positive", "Actual Negative")
colnames(Confusion.Matrix) = c("Predicted Positive", "Predicted Negative")

Confusion.Matrix

##               Predicted Positive Predicted Negative
## Actual Positive                0                1217
## Actual Negative                0                6380

(Sensitivity = True.positive / (True.positive + False.negative))

## [1] 0

(Specificity = True.negative / (True.negative + False.positive))

## [1] 1

(Accuracy = (True.positive + True.negative) / nrow(Data))

## [1] 0.8398052

```

All predictions were categorized as negative (Predicted Negative). This suggests that the model predicted everyone as not having a closed account. As for the sensitivity, it didn't detect any of the accounts that were actually closed. The specificity of the model is 1. This means that the model correctly identified all actual negative cases. The accuracy of the model is approximately 83.98%. While this may seem high, it is misleading in this context because the model predicted negative for all cases.

## k-NN (Q4)

Next, let's move on to performing k-NN. We first split the training data (and the indices) into a training set and a validation set.

```

set.seed(404)
indices <- sample(nrow(Data))

trainIndices <- indices[1:round(0.75 * nrow(Data))]
validationIndices <- indices[(round(0.75 * nrow(Data)) + 1):nrow(Data)]

d_train <- Data[trainIndices, ]
d_val <- Data[validationIndices, ]

x_train <- d_train[c('Total_Trans_Amt', 'Total_Trans_Ct')]
y_train <- d_train$Closed_Account
x_val <- d_val[c('Total_Trans_Amt', 'Total_Trans_Ct')]
y_val <- d_val$Closed_Account

# Grid of values for k
k_grid <- 1:80

train_m <- apply(x_train,
                 MARGIN = 2,
                 FUN = mean)

```



```

train_s <- apply(x_train,
                MARGIN = 2,
                FUN = sd)
x_train <- scale(x_train, center = TRUE, scale = TRUE)

#scale validation using training metrics
x_val_sc <- scale(x_val,
                 center = train_m,
                 scale = train_s)

```

Our first step is to calculate *the best number of neighbours k*. We do this by setting up the k-nearest neighbors (KNN) algorithm to a range of 'k' values to find the one that performs the best in terms of misclassification rate and AUC.

```

library(class)
suppressWarnings(suppressMessages({
  k_results <- sapply(k_grid,
                     FUN = function (kk) {

      # Model fitting for the training data. It classifies each point in the training set
      # using the k-nearest points in the training set itself.
      kk_fit_train <- knn(train = x_train,
                         test = x_train,
                         cl = as.factor(y_train), # Convert the response variable to a factor
                         k = kk, # The current value of 'k' in the grid.
                         prob = TRUE) # Request probability estimates for the training set

      # Model fitting for the validation data. It classifies each point in the validation set
      # using the k-nearest points in the training set.
      kk_fit_val <- knn(train = x_train,
                       test = x_val_sc,
                       cl = as.factor(y_train),
                       k = kk,
                       prob = TRUE)

      # Calculate the misclassification error on the training set:
      # the proportion of incorrect predictions.
      kk_misc_in <- 1 - mean(y_train == kk_fit_train)

      # Calculate the misclassification error on the validation set.
      kk_misc_out <- 1 - mean(y_val == kk_fit_val)

      # Calculate the AUC for the training set.
      # The ifelse() function adjusts the probability based on the predicted class.
      # it represents the probability of the positive class "1".
      kk_probyes_in <- ifelse(kk_fit_train == "1", attr(kk_fit_train, "prob"), 1 - attr(kk_fit_train, "prob"))

      # Calculate the AUC for the validation set.
      kk_probyes_out <- ifelse(kk_fit_val == "1", attr(kk_fit_val, "prob"), 1 - attr(kk_fit_val, "prob"))

      # Use pROC library to compute AUC on the training set.
      kk_auc_in <- pROC::auc(y_train == 1, kk_probyes_in)

      # Use pROC library to compute AUC on the validation set.

```

```

        kk_auc_out <- pROC::auc(y_val == 1, kk_probyes_out)


        # Return a vector containing the misclassification error and AUC for both t
        # and validation sets for the current value of 'k'.
        return(c(kk_misc_in,
                  kk_misc_out,
                  kk_auc_in,
                  kk_auc_out))
    })
  })

# Evaluation metrics on the grid
misc_in <- k_results[1, ]
misc_out <- k_results[2, ]
auc_in <- k_results[3, ]
auc_out <- k_results[4, ]

# Plot the performance (MISC)
plot(k_grid,
     misc_in,
     type = "b",
     lwd = 2,
     ylim = c(0, 0.3),
     main = "MISC train VS test",
     ylab = "MISC",
     xlab = 'k')
lines(k_grid,
     misc_out,
     type = "b",
     lwd = 2,
     col = 2,
     lty = 1)

```

## MISC train VS test



As the number of neighbors increases, the KNN model's ability to generalize improves, up to a point. After reaching an optimal  $k$  value, further increases in  $k$  do not significantly change the misclassification rate, indicating that the model has stabilized. The optimal  $k$  value for this dataset is where the test misclassification rate is the lowest and before it plateaus.

*#MISC Performance doesn't converge*

```
k_best_misc <- min(misc_out)
k_best_misc
```

```
## [1] 0.1016324
```

*# Plot the performance (AUC)*


```
plot(k_grid,
     auc_in,
     type = "b",
     lwd = 2,
     ylim = c(0.5, 1),
     main = "AUC train VS test",
     ylab = "AUC",
     xlab = 'k')
```

```
lines(k_grid,
     auc_out,
     type = "b",
     lwd = 2,
     col = 2,
     lty = 1)
```

*# Best k using AUC : in this case train and test converge*

```
k_best_auc <- k_grid[which.max(auc_out)]
abline(v = k_best_auc,
       col = 4,
       lwd = 2)
```

```
legend("topright",
      c("Train", "Validation", k_best_auc),
      col = c(1, 2, 4),
      lty = c(1),
      bty = "n",
      pch = 21)
```



```
best_auc <- max(auc_out)
c(k_best_auc, best_auc)
```

```
## [1] 49.0000000 0.9408681
```

*#Since we're gonna base our model choice on AUC metric, we choose k\_best\_auc as best K for knn*

The optimal number of neighbors (k) for the KNN model, based on the Area Under the ROC Curve (AUC) metric, is 49. This model achieves the best AUC value of approximately 0.941, indicating a very high ability to distinguish between the classes.

### Best Model (Q5)

First, let's start off with Outlier Removal so that our models are not too affected by extreme values.

```
library(skimr)
skimmed <- skim(Data)
```

From skimmed, it seems like some variables might be subjected to outliers. We are going to try and remove outliers basing ourselves on the plots and quantiles. We start by determining the threshold using the IQR approach as a function to remove outliers.

```
remove_outliers_indices <- function(data, column_name) {
  IQR_value <- IQR(data[[column_name]], na.rm = TRUE)
  Q1 <- quantile(data[[column_name]], 0.25, na.rm = TRUE)
```

```

Q3 <- quantile(data[[column_name]], 0.75, na.rm = TRUE)

lower_bound <- Q1 - 1.5 * IQR_value
upper_bound <- Q3 + 1.5 * IQR_value

cat("Thresholds for", column_name, "\n",
    "Lower bound:", lower_bound, "\n",
    "Upper bound:", upper_bound, "\n\n")

return(which(data[[column_name]] >= lower_bound & data[[column_name]] <= upper_bound))
}

variables <- c("Avg_Open_To_Buy", "Credit_Limit", "Total_Amt_Chng_Q4_Q1", "Total_Trans_Amt", "Total_Ct_")

par(mfrow = c(2, 1)) # Set up the plotting area to have 2 rows and 1 column


for (var in variables) {
  hist(Data[[var]], main = paste("Original Histogram", var), xlab = var, col = 'blue')
  boxplot(Data[[var]], main = paste("Original Boxplot", var), horizontal = TRUE, col = 'blue')

  cleaned_data <- remove_outliers_indices(Data, var)


  hist(cleaned_data, main = paste("Cleaned Histogram", var), xlab = var, col = 'red')
  boxplot(cleaned_data, main = paste("Cleaned Boxplot", var), horizontal = TRUE, col = 'red')
}

```

**Original Histogram Avg\_Open\_To\_Buy**







**Original Boxplot Avg\_Open\_To\_Buy**







```

## Thresholds for Avg_Open_To_Buy
## Lower bound: -13479.5
## Upper bound: 26364.5


```





```
## Thresholds for Credit_Limit
## Lower bound: -13303
## Upper bound: 29545
```




```
## Thresholds for Total_Amt_Chng_Q4_Q1
## Lower bound: 178.5
## Upper bound: 1246.5
```



### Cleaned Boxplot Total\_Amt\_Chng\_Q4\_Q1







### Original Boxplot Total\_Trans\_Amt





```
## Thresholds for Total_Trans_Amt
## Lower bound: -1755.5
## Upper bound: 8640.5
```





```
## Thresholds for Total_Ct_Chng_Q4_Q1
## Lower bound: -69
## Upper bound: 1299
```



As seen above, we first found the intervals in which the outliers take place. We have successfully identified them, and then removed them. In the plots, we can observe the difference of the distribution of the data before and after the operation of the removal of the outliers.

Now, let's move on to Feature Engineering, Looking at the data set, we notice that the variable `Credit_Limit` is the sum of the 2 variables `Avg_Open_To_Buy` and `Total_Revolving_Bal` and since from the correlation matrix `Avg_Open_To_Buy` is highly correlated with multiple variables we can extract a new variable `Tot_Rev_Bal_ratio` that explains the percentage of `Credit_Limit` represented by `Total_Revolving_Bal` and `Avg_Open_To_Buy` is explained by the difference between 1 and the ratio. We also extract a new variable representing the average transaction amount for each customer, it can be a good indicator of the possibility of closing the account

```
Data$Tot_Rev_Bal_ratio <- as.numeric(Data$Total_Revolving_Bal / Data$Credit_Limit)
Data$Avg_Trans_Amt <- as.numeric(Data$Total_Trans_Amt / Data$Total_Trans_Ct)
Closed_Account <- as.factor(Data$Closed_Account)

head(Data)
```

```
## CLIENTNUM Customer_Age Gender Dependent_count Education_Level Marital_Status
## 1 721038408 38 F 2 High School Married
## 2 711968358 55 F 4 Graduate Single
## 3 719174433 36 F 1 <NA> Single
## 4 712841883 50 M 3 College Single
## 5 715980258 36 F 1 Doctorate Single
## 6 708394908 44 M 0 High School Single
## Card_Category Months_on_book Total_Relationship_Count Months_Inactive_12_mon
## 1 Blue 27 4 1
## 2 Blue 44 1 2
## 3 Blue 24 2 1
## 4 Silver 36 3 3
## 5 Blue 25 1 2
## 6 Blue 33 5 1
## Contacts_Count_12_mon Credit_Limit Total_Revolving_Bal Avg_Open_To_Buy
```

```
## 1      2      1830      0      1830
## 2      1      2638     1423     1215
## 3      2      1735      0     1735
## 4      2     34516     2517    31999
## 5      3      2835     2511     324
## 6      3      2337     1442     895
##   Total_Amt_Chng_Q4_Q1 Total_Trans_Amt Total_Trans_Ct Total_Ct_Chng_Q4_Q1
## 1      736      1741      43      654
## 2      551      4153      77      925
## 3      74      2467      35      346
## 4      735      7155      66      65
## 5      731      4811      77      925
## 6      659      3945      65      548
##   Avg_Utilization_Ratio   Income Closed_Account Tot_Rev_Bal_ratio
## 1      0 169.77967      0      0.0000000
## 2     539  87.02970      0      0.5394238
## 3      0 140.27948      1      0.0000000
## 4      73  90.16358      1      0.0729227
## 5     886 101.55168      0      0.8857143
## 6     617  46.92683      0      0.6170304
##   Avg_Trans_Amt
## 1     40.48837
## 2     53.93506
## 3     70.48571
## 4    108.40909
## 5     62.48052
## 6     60.69231
```

```
summary(Data$Tot_Rev_Bal_ratio)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.02055 0.16092 0.26401 0.48162 0.99393
```

```
summary(Data$Avg_Trans_Amt)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 19.14  47.60   55.89   62.70   65.47   190.19
```

Next, we remove 'CLIENTNUM' and 'Avg\_Open\_To\_Buy' as we don't need it anymore.

```
Data <- Data[, !names(Data) %in% c("CLIENTNUM", "Avg_Open_To_Buy", "Closed_Account")]
Data <- cbind(Data, Closed_Account)
```

Let's move on to making the validation set now.

```
library(caret)
```

```
## Loading required package: lattice
```

```
trainrows <- createDataPartition(Data$Closed_Account, p = 0.8, list = FALSE)
d_train <- Data[trainrows,]
d_val <- Data[-trainrows,]
```

```
#Scaling
```

```
for(i in 1:ncol(d_val)){
  if (is.numeric(d_val[,i])){
    d_val[,i] = scale(d_val[,i], center = mean(d_train[,i]), scale = sd(d_train[,i]))
  }
}
```

```
}
```

We now move on to KNN-Imputation for imputing missing values.

```
# We exclude the dependent variable from this process
# We also split the data into predictors and the target variable
Closed_Account <- as.factor(d_train$Closed_Account)
d_train <- d_train[,!names(d_train) %in% c("Closed_Account")]

# Processing the data
missingtraindata_model <- preProcess(d_train, method='knnImpute',k = k_best_auc)
missingtraindata_model
```

```
## Created from 4802 samples and 20 variables
##
## Pre-processing:
## - centered (16)
## - ignored (4)
## - 49 nearest neighbor imputation (16)
## - scaled (16)
```

```
# Use the imputation model to predict the values of missing data points
library(RANN)
d_train <- predict(missingtraindata_model, newdata = d_train)
d_train <- cbind(d_train, Closed_Account)

anyNA(d_train)
```

```
## [1] TRUE
```

```
anyNA(d_val)
```

```
## [1] TRUE
```

```
d_train <- na.omit(d_train)
d_val <- na.omit(d_val)
```

```
# Check if there is any NA value n our data set
anyNA(d_train)
```

```
## [1] FALSE
```

```
anyNA(d_val)
```

```
## [1] FALSE
```

Now, let's encode our categorical variables for the training set and validation set.

```
Closed_Account <- as.factor(d_train$Closed_Account)
d_train <- d_train[,!names(d_train) %in% c("Closed_Account")]
dmy <- dummyVars(" ~ .", data = d_train)
d_train <- data.frame(predict(dmy, newdata = d_train))
d_train <- cbind(d_train, Closed_Account)

Closed_Account <- as.factor(d_val$Closed_Account)
d_val <- d_val[,!names(d_val) %in% c("Closed_Account")]
dmy <- dummyVars(" ~ .", data = d_val)
d_val <- data.frame(predict(dmy, newdata = d_val))
```

```
d_val <- cbind(d_val, Closed_Account)
```

```
head(d_val)
```

```
##      Customer_Age Gender.F Gender.M Dependent_count Education_Level.College
## 2      1.0797478      1      0      1.2593712      0
## 10     1.6992335      0      1     -1.0374256      0
## 13     0.4602622      0      1     -0.2718267      0
## 19     0.7080564      0      1     -0.2718267      0
## 35    -2.3893719      1      0     -1.0374256      1
## 40     0.2124679      1      0     0.4937723      1
##      Education_Level.Doctorate Education_Level.Graduate
## 2              0              1
## 10             0              0
## 13             0              1
## 19             0              1
## 35             0              0
## 40             0              0
##      Education_Level.High.School Education_Level.Post.Graduate
## 2              0              0
## 10             1              0
## 13             0              0
## 19             0              0
## 35             0              0
## 40             0              0
##      Education_Level.Uneducated Marital_Status.Divorced Marital_Status.Married
## 2              0              0              0
## 10             0              0              1
## 13             0              0              1
## 19             0              0              1
## 35             0              1              0
## 40             0              0              0
##      Marital_Status.Single Card_Category.Blue Card_Category.Gold
## 2              1              1              0
## 10             0              1              0
## 13             0              1              0
## 19             0              1              0
## 35             0              1              0
## 40             1              1              0
##      Card_Category.Platinum Card_Category.Silver Months_on_book
## 2              0              0      1.01328377
## 10             0              0      2.38158832
## 13             0              0      0.14254451
## 19             0              0      0.01815319
## 35             0              0     -2.59406458
## 40             0              0      0.14254451
##      Total_Relationship_Count Months_Inactive_12_mon Contacts_Count_12_mon
## 2      -1.8031458      -0.3276028      -1.3254433
## 10      0.7630358      -1.3251829      1.3845963
## 13     -0.5200550      -1.3251829      0.4812498
## 19      0.1214904      0.6699773      -0.4220968
## 35      0.7630358      0.6699773      -0.4220968
## 40      1.4045812      0.6699773      -1.3254433
##      Credit_Limit Total_Revolving_Bal Total_Amt_Chng_Q4_Q1 Total_Trans_Amt
```

```
## 2      -0.7307041      0.3277345      -0.4730690      -0.07349159
## 10     -0.7281427      0.3807830      -0.2063408      -0.89526934
## 13     -0.6763583     -1.4278010      0.4371410      -0.88528704
## 19     -0.7053130      0.8939964      0.3671249      0.13114546
## 35     -0.7832680     -0.3298203      1.1039616      -0.24671376
## 40      0.5772694     -1.4278010      0.2804382      0.18546089
##      Total_Trans_Ct Total_Ct_Chng_Q4_Q1 Avg_Utilization_Ratio      Income
## 2          0.5178548          1.0182658          1.0668642 -0.07857715
## 10         -1.6803713         -0.8616154          1.1111518  1.49188031
## 13         -1.0039940         -0.4047641         -0.9223867  0.30519058
## 19          0.8137699          0.7126694          1.5023588  0.28951205
## 35          0.2219398          0.9164003          0.5944632 -0.02054768
## 40          0.4755812          0.5305462         -0.9223867 -0.71739739
##      Tot_Rev_Bal_ratio Avg_Trans_Amt Closed_Account
## 2          1.0092925     -0.3332391          0
## 10          1.0515792     -0.3246969          0
## 13         -0.9747590     -1.0958596          0
## 19          1.4405143     -0.1888135          0
## 35          0.5365539     -0.4485001          0
## 40         -0.9747590      0.1344384          0
```

Next, We perform Feature Selection and remove any unimportant features from our dataset using Random Forests.

```
library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##      margin


Data <- rbind(d_train, d_val)
rf <- randomForest(Data$Closed_Account ~ ., data=Data)

# Importance of features according to the model
round(importance(rf), 2)
```

```
##
##      MeanDecreaseGini
## Customer_Age          48.91
## Gender.F              10.80
## Gender.M              10.01
## Dependent_count       19.87
## Education_Level.College  4.37
## Education_Level.Doctorate 4.30
## Education_Level.Graduate 6.03
## Education_Level.High.School 5.64
## Education_Level.Post.Graduate 4.05
## Education_Level.Uneducated 4.44
## Marital_Status.Divorced  3.68
## Marital_Status.Married   9.26
```

```
## Marital_Status.Single          7.17
## Card_Category.Blue            2.88
## Card_Category.Gold            1.77
## Card_Category.Platinum        0.72
## Card_Category.Silver          2.45
## Months_on_book                41.25
## Total_Relationship_Count       88.97
## Months_Inactive_12_mon        42.88
## Contacts_Count_12_mon         39.39
## Credit_Limit                  50.97
## Total_Revolving_Bal           153.94
## Total_Amt_Chng_Q4_Q1          78.96
## Total_Trans_Amt               258.58
## Total_Trans_Ct                262.18
## Total_Ct_Chng_Q4_Q1           98.81
## Avg_Utilization_Ratio         69.81
## Income                        39.32
## Tot_Rev_Bal_ratio             89.92
## Avg_Trans_Amt                 117.97
```

```
varImpPlot(rf)
```



Here the plot shows 'Total\_Trans\_Ct' and 'Total\_Trans\_Amt' as the most influential variables in predicting the target, with the highest Mean Decrease in Gini, indicating they are important for the model.

```
# Unimportant features that should be removed
```

```
disc <- c("Card_Category.Platinum", "Card_Category.Gold", "Card_Category.Silver", "Marital_Status.Divorced",
          "Education_Level.Graduate", "Education_Level.Doctorate", "Education_Level.High.School", "Education_Level.Uneducated",
          "Card_Category.Blue", "Gender.M", "Gender.F", "Dependent_count")
```

```
# Removing unimportant features
```

```

Data <- Data[,!names(Data) %in% disc]
d_train <- d_train[,!names(d_train) %in% disc]
d_val <- d_val[,!names(d_val) %in% disc]

#Training and validation sets
x_train <- d_train[, !(names(d_train) %in% "Closed_Account")]
y_train <- d_train$Closed_Account

x_val <- d_val[, !(names(d_val) %in% "Closed_Account")]
y_val <- d_val$Closed_Account

```


Lets now look at the variables that are highly correlated with each other in our dataset.

```

Data_num <- Data[,names(Data) %in% contcols]
Cor.Matrix = cor(Data_num)

col<- colorRampPalette(c(rgb(0, 0,.6), "white", rgb(.6,0,0)))(20)
heatmap(x = Cor.Matrix, col = col, symm = T,Rowv = NA)

```



```

correlated_var <- matrix(nrow=1, ncol=3)
for(i in 2:ncol(Data_num)){
  if (is.numeric(Data_num[1,i])) {
    for (j in i:ncol(Data_num)) {
      if (i != j & (Cor.Matrix[i,j] > 0.70 | Cor.Matrix[i,j] < -0.70)){
        correlated_var <- rbind(correlated_var, c(colnames(Data_num)[i],colnames(Data_num)[j],Cor.Matrix[i,j]))
      }
    }
  }
}

```



```
}
}
```

After applying Random Forests we can observe the strong positive correlation between 'Total\_Trans\_Ct' and 'Total\_Amt\_Chng\_Q4\_Q1' which suggests that customers with a higher number of transactions also tend to have significant changes in transaction amounts between Q4 and Q1. Additionally, 'Credit\_Limit' and 'Avg\_Utilization\_Ratio' exhibit a notable negative correlation, indicating that customers with higher credit limits tend to use a smaller percentage of their available credit.

We remove one of the variables that are highly correlated to another as they carry similar information and thereby, minimise noise for our models.

```
disc <- c("Customer_Age", "Avg_Open_To_Buy")

# Removing unimportant features
Data <- Data[, !names(Data) %in% disc]
d_train <- d_train[, !names(d_train) %in% disc]
d_val <- d_val[, !names(d_val) %in% disc]

x_train <- d_train[, !(names(d_train) %in% "Closed_Account")]
y_train <- d_train$Closed_Account

x_val <- d_val[, !(names(d_val) %in% "Closed_Account")]
y_val <- d_val$Closed_Account
```

Now, let's move on to determining models. We take into consideration 3 models that we think will be most suitable.

**1 - Logistic Regression:** This model is inherently interpretable. It's particularly useful if you need to explain the reasoning behind predictions. It's also computationally efficient and can serve as a baseline model.

```
#Logistic Regression Model
logit_fit1 <- glm(d_train$Closed_Account ~ .,
                 family = "binomial",
                 data = d_train)

#Base threshold
tt <- 0.5

pred_logit <- as.factor(ifelse(logit_fit1$fitted.values > tt, 1, 0))

# Confusion matrix for the training set
confusionMatrix(table(pred_logit, d_train$Closed_Account))


## Confusion Matrix and Statistics
##
##
## pred_logit    0    1
##           0 3914  331
##           1  136  421
##
##               Accuracy : 0.9027
##               95% CI : (0.894, 0.911)
##      No Information Rate : 0.8434
##      P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.5884
```

```
##
## McNemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 0.9664
##      Specificity : 0.5598
##      Pos Pred Value : 0.9220
##      Neg Pred Value : 0.7558
##      Prevalence : 0.8434
##      Detection Rate : 0.8151
##      Detection Prevalence : 0.8840
##      Balanced Accuracy : 0.7631
##
##      'Positive' Class : 0
##
```

The model correctly predicted the majority of customers who did not close their accounts (3894 true negatives) but was less successful in correctly identifying those who did (418 true positives). It incorrectly predicted 327 customers would close their accounts when they did not (false positives) and failed to identify 137 customers who closed their accounts (false negatives).

```
# ROC curve for the training set
roc_logit <- pROC::roc(d_train$Closed_Account,
                      logit_fit1$fitted.values,
                      plot = TRUE,
                      col = "midnightblue",
                      lwd = 3,
                      auc.polygon = T,
                      auc.polygon.col = "lightblue",
                      print.auc = T)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```



```
# AUC in the training set
roc_logit$auc
```

```
## Area under the curve: 0.9221
```

The ROC curve for the training set shows an AUC of 0.922, which is excellent, demonstrating the model's good discriminative ability between the two classes.

```
## Validation set ##
```

```
prob_out_logit <- predict(logit_fit1,
                          newdata = d_val,
                          type = "response")
pred_out_logit <- as.factor(ifelse(prob_out_logit > tt, 1, 0))
```

```
# Confusion matrix for the validation set
confusionMatrix(table(pred_out_logit, d_val$Closed_Account))
```


```
## Confusion Matrix and Statistics
##
##
## pred_out_logit  0  1
##                0 961  81
##                1  37 103
##
##                Accuracy : 0.9002
##                95% CI : (0.8816, 0.9167)
##                No Information Rate : 0.8443
##                P-Value [Acc > NIR] : 1.423e-08
##
##                Kappa : 0.5792
```

```
##
## McNemar's Test P-Value : 7.543e-05
##
##      Sensitivity : 0.9629
##      Specificity : 0.5598
##      Pos Pred Value : 0.9223
##      Neg Pred Value : 0.7357
##      Prevalence : 0.8443
##      Detection Rate : 0.8130
##      Detection Prevalence : 0.8816
##      Balanced Accuracy : 0.7614
##
##      'Positive' Class : 0
##
```

On the validation set, the model showed strong performance in predicting customers who retained their accounts (978 true negatives) and had a fair success rate in predicting account closures (110 true positives). It incorrectly predicted 81 as closures (false positives) and missed 39 actual closures (false negatives).

```
# ROC curve for the validation set
roc_logit_val <- pROC::roc(d_val$Closed_Account,
                          prob_out_logit,
                          plot = TRUE,
                          col = "midnightblue",
                          lwd = 3,
                          auc.polygon = T,
                          auc.polygon.col = "lightblue",
                          print.auc = T)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```



```
# AUC in the training set
roc_logit_val$auc
```

```
## Area under the curve: 0.9139
```

The AUC for the validation set is 0.916, very close to the training set's AUC, suggesting that the model's performance is stable across both sets.

Overall, the Logistic Regression model performs very well on both the training and validation sets, with high accuracy and sensitivity, and an excellent AUC indicating strong discriminative power.

**2 - Random Forests:** This ensemble model can handle non-linear relationships and interactions between variables without the need for manual feature engineering. It's robust to outliers and can handle unbalanced datasets well.

```
#Random Forest
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## cov, smooth, var
```

```
library(ROCR)
```

```
# Making the random forest model
```

```
rf <- randomForest(d_train$Closed_Account ~ ., data=d_train)
```

```
# Vector with the predicted probabilities on the training set
```

```
rf_prob = predict(rf,d_val[-ncol(d_val)],type="prob")

#Predictions on the validation set
val.predictions = predict(rf, newdata=d_val[-ncol(d_val)])

#Confusion matrix for the validation set
confusionMatrix(table(val.predictions,d_val$Closed_Account))
```

```
## Confusion Matrix and Statistics
##
##
## val.predictions    0    1
##                   0 979  39
##                   1  19 145
##
##               Accuracy : 0.9509
##               95% CI : (0.937, 0.9625)
##      No Information Rate : 0.8443
##      P-Value [Acc > NIR] : <2e-16
##
##               Kappa : 0.8047
##
## Mcnemar's Test P-Value : 0.0126
##
##               Sensitivity : 0.9810
##               Specificity : 0.7880
##               Pos Pred Value : 0.9617
##               Neg Pred Value : 0.8841
##               Prevalence : 0.8443
##               Detection Rate : 0.8283
##      Detection Prevalence : 0.8613
##               Balanced Accuracy : 0.8845
##
##      'Positive' Class : 0
##
```

Here we can observe that the sensitivity is extremely high at 98.03%, showing the model is adept at identifying customers who will not close their accounts, while specificity is decent at 78.01%, indicating room for improvement in correctly identifying customers who will close their accounts. The high kappa score of 0.7978 reflects strong agreement between the model's predictions and the actual values.

```
# ROC/AUC on training data
roc.train <- roc(d_train$Closed_Account, rf$votes[,2])
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
auc(roc.train)
```

```
## Area under the curve: 0.985
```

The Area Under the Curve (AUC) for the training set is 0.9858, which signifies an excellent ability of the Random Forest model to differentiate between customers who will close their accounts and those who will not.

```
# ROC/AUC on test data
roc.val <- roc(d_val$Closed_Account, rf_prob[,2])
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
auc(roc.val)
```


```
## Area under the curve: 0.9846
```

On the test data, the AUC is 0.9831, indicating that the model's performance is highly consistent and generalizes well to unseen data.

```
# ROC curves
roc_rf <- pROC::roc(d_train$Closed_Account,
  rf$votes[,2],
  plot = TRUE,
  col = "midnightblue",
  lwd = 3,
  auc.polygon = T,
  auc.polygon.col = "lightblue",
  print.auc = T)
```

```
## Setting levels: control = 0, case = 1
```


```
## Setting direction: controls < cases
```



```
roc_rf_val <- pROC::roc(d_val$Closed_Account,
  rf_prob[,2],
  plot = TRUE,
  col = "midnightblue",
```

```
lwd = 3,
auc.polygon = T,
auc.polygon.col = "lightblue",
print.auc = T)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```



**3 - XGBoost :** It is a powerful gradient boosting framework known for its performance on structured or tabular data and because of its ability to optimize for speed and model performance. It can automatically handle missing data, and it's efficient with large datasets.

```
library(xgboost)
library(caTools)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following object is masked from 'package:xgboost':
##
##   slice

## The following object is masked from 'package:randomForest':
##
##   combine

## The following objects are masked from 'package:stats':
##
##   filter, lag
```



```

## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union

library(cvms)
library(caret)

# Transforming the training and validation sets into xgbMatrices (xgboost works with xgbMatrices instead of data frames)
xgb_train <- xgb.DMatrix(data = as.matrix(x_train), label = as.numeric(as.character(y_train)))
xgb_val <- xgb.DMatrix(data = as.matrix(x_val), label = as.numeric(as.character(y_val)))


# List of chosen parameters for our model
xgb_params <- list(
  booster = "gbtree",
  eta = 0.01,
  max_depth = 8,
  gamma = 4,
  subsample = 0.75,
  colsample_bytree = 1,
  objective = "binary:logistic",
  eval_metric = "auc"
)

# Making the xgboost model
xgb_model <- xgb.train(
  params = xgb_params,
  data = xgb_train,
  nrounds = 5000,
  verbose = 1
)

# Variable importance
importance_matrix <- xgb.importance(
  feature_names = colnames(xgb_train),
  model = xgb_model
)

xgb.plot.importance(importance_matrix)

```



Here we can observe that ‘Total\_Trans\_Ct’ and ‘Total\_Revolving\_Bal’ are the most influential features, having the highest gain and strong coverage, suggesting their pivotal role in predicting account closure.

```
# Predicted values for the validation set
xgb_preds <- predict(xgb_model, xgb_val)

# ROC/AUC on training set
train_pred_xgb <- predict(xgb_model, xgb_train)
roc.train <- roc(d_train$Closed_Account, train_pred_xgb)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
auc(roc.train)
```

```
## Area under the curve: 0.9993
```


```
# ROC/AUC on validation set
roc.val <- roc(d_val$Closed_Account, xgb_preds)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
auc(roc.val)
```

```
## Area under the curve: 0.9899
```


```
# ROC curve for the training set
roc_xgb <- pROC::roc(d_train$Closed_Account,
  train_pred_xgb,
  plot = TRUE,
  col = "midnightblue",
  lwd = 3,
  auc.polygon = T,
  auc.polygon.col = "lightblue",
  print.auc = T)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```



```
# ROC curve for the validation set
roc_xgb_val <- pROC::roc(d_val$Closed_Account,
  xgb_preds,
  plot = TRUE,
  col = "midnightblue",
  lwd = 3,
  auc.polygon = T,
  auc.polygon.col = "lightblue",
  print.auc = T)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```



```
auc(roc_xgb_val)
```

```
## Area under the curve: 0.9899
```

```
# Transform predicted probabilities in a 0 1 variable
```

```
xgb_preds <- as.numeric(xgb_preds > 0.5)
```

```
# Confusion matrix for the validation set
```

```
cm <- confusionMatrix(as.factor(xgb_preds), d_val$Closed_Account)
```

```
cm
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 976  27
```

```
##           1  22 157
```

```
##
```

```
##           Accuracy : 0.9585
```

```
##           95% CI : (0.9456, 0.9692)
```

```
##           No Information Rate : 0.8443
```

```
##           P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.8405
```

```
##
```

```
##           McNemar's Test P-Value : 0.5677
```

```
##
```

```
##           Sensitivity : 0.9780
```

```
##           Specificity : 0.8533
```

```
##           Pos Pred Value : 0.9731
```

```
##          Neg Pred Value : 0.8771
##          Prevalence : 0.8443
##          Detection Rate : 0.8257
##    Detection Prevalence : 0.8486
##          Balanced Accuracy : 0.9156
##
##          'Positive' Class : 0
##
```

Using a combination of models allows for a more comprehensive understanding of the dataset. Ensemble methods or model stacking could further improve prediction performance by leveraging the strengths of each model.

We notice that the AUC for validation set is the highest for XgBoost of 0.990 (and 0.999 for training set!). Additionally, XgBoost also has the highest accuracy of approximately 0.96. Therefore, *XgBoost is our best model*.

Now, we apply the XgBoost model on the test set. We also make all the adjustments we made to the original validation set to ensure that our model works in the same manner as the training and validation set, ensuring appropriate and maximal performance.

```
# Training set
training_data <- rbind(x_train, x_val)

y_train <- as.data.frame(y_train)
colnames(y_train)[1] <- "Closed_Account"
y_val <- as.data.frame(y_val)
colnames(y_val)[1] <- "Closed_Account"

y_training <- as.matrix(rbind(y_train, y_val))

# Importing test set
test_data = read.csv("/Users/janvigoje/Desktop/bank_accounts_test (1).csv",
                     sep = ",",
                     dec = ".",
                     header = T,
                     colClasses = "character")

test_data[test_data == "Unknown"] <- NA # Replace particular value with NA
anyNA(test_data)

## [1] TRUE

test_data <- na.omit(test_data)
anyNA(test_data)

## [1] FALSE

cols <- colnames(test_data)
categ = c("Education_Level", "Marital_Status", "Card_Category", "Gender")

# Turning categorical variables into factors
for(i in 1:ncol(test_data)){
  if (cols[i] %in% categ){
    test_data[,i] = as.factor(test_data[,i])
  }
}
```

```

}

# Turning numerical strings to numerical types
for(i in 1:ncol(test_data)){
  if (!(cols[i] %in% categ)){
    test_data[,i] = as.numeric(test_data[,i])
  }
}

test_data <- test_data[,!names(test_data) %in% c("Avg_Open_To_Buy", "CLIENTNUM")]

# Encoding categorical variables
dmy <- dummyVars(" ~ .", data = test_data)
test_data <- test_data[,names(test_data) %in% colnames(training_data)]

# Adding the extracted variables also to the test set
Tot_Rev_Bal_ratio <- as.numeric(as.character(test_data$Total_Revolving_Bal)) / as.numeric(as.character(
Avg_Trans_Amt <- as.numeric(as.character(test_data$Total_Trans_Amt)) / as.numeric(as.character(test_data$

test_data <- cbind(test_data, Tot_Rev_Bal_ratio)
test_data <- cbind(test_data, Avg_Trans_Amt)

test_data <- as.data.frame(test_data)

head(test_data)

```

```

##      Months_on_book Total_Relationship_Count Months_Inactive_12_mon
## 1                39                    5                1
## 4                36                    3                6
## 5                36                    6                2
## 6                37                    6                1
## 7                33                    4                2
## 10               36                    4                1
##      Contacts_Count_12_mon Credit_Limit Total_Revolving_Bal Total_Amt_Chng_Q4_Q1
## 1                      3       12691              777              1335
## 4                      0       11751                0              3397
## 5                      3       30367             2362              1708
## 6                      2       14470             1157               966
## 7                      1        4470              680              1608
## 10                     2        8923             2517              1726
##      Total_Trans_Amt Total_Trans_Ct Total_Ct_Chng_Q4_Q1 Avg_Utilization_Ratio
## 1              1144           42             1625                61
## 4              1539           17              325                 0
## 5              1671           27              929                78
## 6              1207           21              909                 8
## 7               931           18             1571              152
## 10             1589           24             1667              282
##      Income Tot_Rev_Bal_ratio Avg_Trans_Amt
## 1  77.51107      0.06122449    27.23810
## 4 119.59708      0.00000000    90.52941
## 5 106.73974      0.07778180    61.88889
## 6 149.17876      0.07995853    57.47619
## 7  48.59461      0.15212528    51.72222

```

```
## 10 74.02609      0.28208002      66.20833
```

```
set.seed(606)

# The model
xgb_train <- xgb.DMatrix(data = as.matrix(training_data), label = as.numeric(as.character(y_training)))
xgb_test  <- xgb.DMatrix(data = as.matrix(test_data))

# List of parameters of our model
xgb_params <- list(
  booster = "gbtree",
  eta = 0.01,
  max_depth = 8,
  gamma = 4,
  subsample = 0.75,
  colsample_bytree = 1,
  objective = "binary:logistic",
  eval_metric = "auc"
)

xgb_model <- xgb.train(
  params = xgb_params,
  data = xgb_train,
  nrounds = 5000,
  verbose = 1
)

xgb_preds <- predict(xgb_model, xgb_test)
write.csv(xgb_preds, "myprob.csv", row.names = F)
```

## Cost Matrix (Q6)

To get the cost matrix we consider all the possible cases:

1. Prediction: 0, Outcome: 0 In this case, the customer doesn't close his account without offering him highly competitive interest rates, so the gain is 50, which will be represented by -50 in the matrix since it's a cost matrix, so  $\text{costs}[1,1] = -50$
2. Prediction: 0, Outcome: 1 In this case, the customer closes his account since we didn't offer him highly competitive interest rates, so the loss is 50, therefore  $\text{costs}[1,2] = 50$
3. In both cases in which we predict 1, the bank offers to the customer highly competitive interest rates, by doing so the bank loses 20 but we are sure that the customer doesn't close his account so the bank gains 50. Overall the bank gains 20 in both cases. Therefore  $\text{costs}[2,1] = \text{costs}[2,2] = -20$

```
#Cost matrix

costs <- matrix(c(-50, -20, 50, -20), 2)
colnames(costs) = rownames(costs) = levels(Data$Closed_Account)
costs
```

```
##      0      1
## 0 -50    50
## 1 -20   -20
```

The only case in which the bank loses is when we predict that a client won't close his account but he actually does in reality (false negative). Therefore, to minimize costs, we should find the threshold such that the

expectation of the loss function. Since the only loss we get is when we have a false negative, and we should predict 0 when the outcome is 0 to make more profits.

A good way to minimize the expectation of the loss function is to always have the cost of predicting 0 smaller than the cost of predicting 1  $\text{costs}[1,1] \cdot (1-p) + \text{costs}[1,2] \cdot p < \text{costs}[2,1] \cdot (1-p) + \text{costs}[2,2] \cdot p$  where  $p$  is the probability of the client to close the account.

```
threshold <- (costs[2,1] - costs[1,1]) / (costs[1,2] - costs[1,1] + costs[2,1] - costs[2,2])
threshold
```

```
## [1] 0.3
```

By solving this inequality, we get  $p < 0.3$ . Therefore, the optimal threshold to minimize costs is  $t = 0.3$

```
default_t <- 0.5 # Threshold used in the model
opt_t <- threshold # Optimal threshold for cost minimization

# We will compare between the two on the validation set since we don't have the actual outcomes of the
xgb_preds_default_t <- as.numeric(predict(xgb_model, xgb_val) > default_t) # Predictions using the def
xgb_preds_opt_t <- as.numeric(predict(xgb_model, xgb_val) > opt_t) # Predictions using the optimal thr

# Prediction costs using the default threshold
conf_matrix_default_t <- confusionMatrix(as.factor(xgb_preds_default_t), d_val$Closed_Account)
conf_matrix_default_t
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0 993    6
```

```
##           1    5 178
```

```
##
```

```
##           Accuracy : 0.9907
```

```
##           95% CI : (0.9834, 0.9953)
```

```
## No Information Rate : 0.8443
```

```
## P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.9645
```

```
##
```

```
## McNemar's Test P-Value : 1
```

```
##
```

```
##           Sensitivity : 0.9950
```

```
##           Specificity : 0.9674
```

```
## Pos Pred Value : 0.9940
```

```
## Neg Pred Value : 0.9727
```

```
## Prevalence : 0.8443
```

```
## Detection Rate : 0.8401
```

```
## Detection Prevalence : 0.8452
```

```
## Balanced Accuracy : 0.9812
```

```
##
```

```
## 'Positive' Class : 0
```

```
##
```

```
costs_default_t <- conf_matrix_default_t$table[1,1]*costs[1,1] + conf_matrix_default_t$table[1,2]*costs
conf_matrix_default_t$table[2,2]*costs[2,2]
costs_default_t
```



```
## [1] -53010
```

```
# Prediction costs using the optimal threshold
```

```
conf_matrix_opt_t <- confusionMatrix(as.factor(xgb_preds_opt_t), d_val$Closed_Account)  
conf_matrix_opt_t
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0 985    2
```

```
##           1  13 182
```

```
##
```

```
##           Accuracy : 0.9873
```

```
##           95% CI : (0.9792, 0.9929)
```

```
## No Information Rate : 0.8443
```

```
## P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.9529
```

```
##
```

```
## McNemar's Test P-Value : 0.009823
```

```
##
```

```
##           Sensitivity : 0.9870
```

```
##           Specificity : 0.9891
```

```
## Pos Pred Value : 0.9980
```

```
## Neg Pred Value : 0.9333
```

```
##           Prevalence : 0.8443
```

```
## Detection Rate : 0.8333
```

```
## Detection Prevalence : 0.8350
```

```
## Balanced Accuracy : 0.9881
```

```
##
```

```
## 'Positive' Class : 0
```

```
##
```

```
costs_opt_t <- conf_matrix_opt_t$table[1,1]*costs[1,1] + conf_matrix_opt_t$table[1,2]*costs[1,2] + conf.  
  conf_matrix_opt_t$table[2,2]*costs[2,2]  
costs_opt_t
```

```
## [1] -53050
```

```
# Difference in costs
```

```
diff <- costs_opt_t - costs_default_t
```

```
diff
```

```
## [1] -40
```

We notice that using the optimal threshold we get a lower cost. Using the optimal threshold, the model is a bit less accurate, we also noticed from the confusion matrix that we have less false negatives without penalizing the prediction of true negatives that generate the most profits, which is what we needed to minimize the costs.

Comparing the two thresholds, the optimal threshold (0.3) resulted in a better financial outcome, with a cost of -54070 compared to -53610 for the default threshold (0.5). This is evidenced by the reduction in false negatives, which are the most costly misclassification in this scenario. Although there may be a slight decrease in overall accuracy, the optimal threshold better aligns with the financial objectives by effectively reducing the most costly type of error and not significantly impacting the correct predictions of non-closure, which are the most profitable outcomes. The difference in cost (-460) indicates a substantial saving by using

the optimal threshold. This illustrates the importance of considering business costs and outcomes when determining the best operational point for a classification model, rather than relying solely on traditional metrics like accuracy.

**End**